



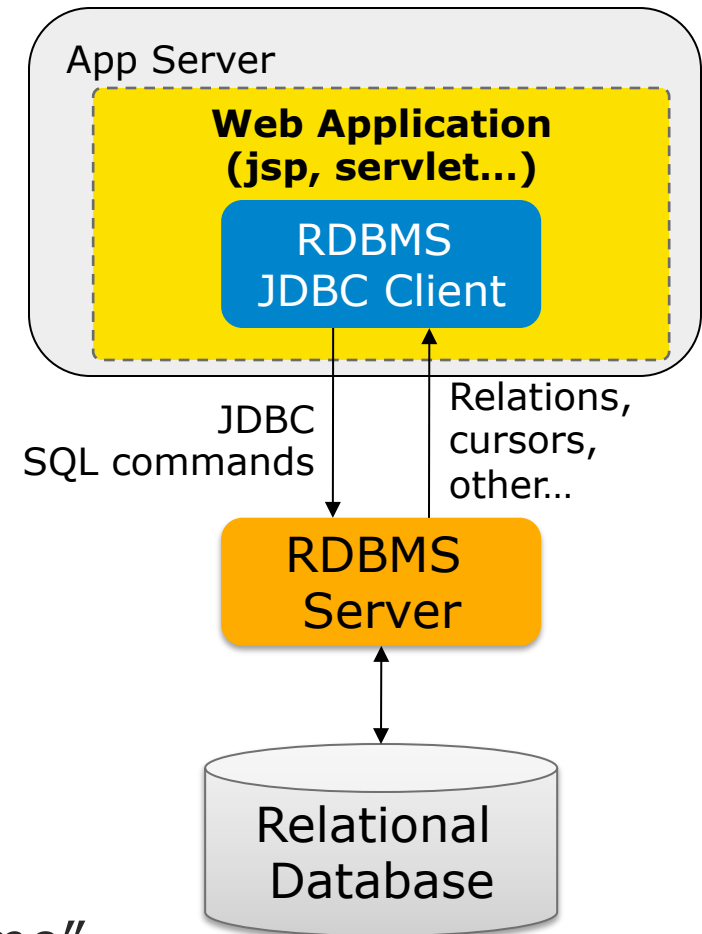
# **MAS 201**

## **Database design & SQL programming**



# Applications' View of a Relational Database Management System (RDBMS): Why use it?

- Persistent data structure
  - Large volume of data
- **High-level language/API for reading (querying) & writing (inserting, deleting, updating)**
  - Automatically optimized
- Transaction management (ACID)
  - Atomicity: all or none happens, despite failures & errors
  - Consistency
  - Isolation: appearance of “one at a time”
  - Durability: recovery from failures and other errors



# OLTP Vs OLAP use cases

## OLTP

- Support quick ACID transactions
- Eg, Bank application that manages transactions

## OLAP

- Perform analytics on the database
- Eg, Bank application analyzing customer profiles towards marketing

- All well-known databases can do both
- But may not be very efficient in analytics
- Many new databases focused on analytics
  - Organizations may have two databases – OLTP vs OLAP
  - Or 3+
- The jury is out on whether two kinds of databases will be needed

# Data Structure: Relational Model

- **Relational Databases:**  
Schema + Data
- **Schema:**
  - collection of *tables* (also called *relations*)
  - each table has a set of *attributes* (aka *columns*)
  - no repeating table names, no repeating attributes in one table
- **Data** (also called *instance*):
  - set of *tuples* (aka *rows*)
  - tuples have one atomic *value* for each attribute

Movie		
ID	Title	Actor
1	Wild	Winger
2	Sky	Winger
3	Reds	Beatty
4	Tango	Brando
5	Tango	Winger
7	Tango	Snyder

Schedule		
ID	Theater	Movie
1	Odeon	1
2	Forum	3
3	Forum	2

# Data Structure: Primary Keys; Foreign Keys are value-based pointers

**Schedule**

ID	Theater	Movie
1	Odeon	1
2	Forum	3
3	Forum	2

**Movie**

ID	Title	Director	Actor
1	Wild	Lynch	Winger
2	Sky	Berto	Winger
3	Reds	Beatty	Beatty
4	Tango	Berto	Brando
5	Tango	Berto	Winger
7	Tango	Berto	Snyder

- “ID is *primary key* of **Schedule**” => its value is unique in **Schedule.ID**
- “**Schedule.Movie** is foreign key (referring) to **Movie.ID**” means every **Movie** value of **Schedule** also appears as **Movie.ID**
- Intuitively, **Schedule.Movie** operates as pointer to **Movie(s)**

# Schema design has its own intricacies

Schedule			Movie			
ID	Theater	Movie	ID	Title	Director	Actor
1	Odeon	1	1	Wild	Lynch	Winger
2	Forum	3	2	Sky	Berto	Winger
3	Forum	2	3	Reds	Beatty	Beatty
			4	Tango	Berto	Brando
			5	Tango	Berto	Winger
			7	Tango	Berto	Snyder

- This example is a bad schema design!
- Problems
  - Change the name of a theater
  - Change the name of a movie's director
  - What about theaters that play no movie?

# How to Design a Database and Avoid Bad Decisions

- With experience...
- Normalization rules of database design instruct how to turn a “bad” design into a “good” one
  - a well-developed mathematical theory
  - no guidance on how to start
  - does not solve all problems
- MAS 201: Think **entities and relationships** – then translate them to tables
- MAS 201: The special case of star & snowflake schemas

# **Designing Schemas Using Entity-Relationship modeling**

The Basics



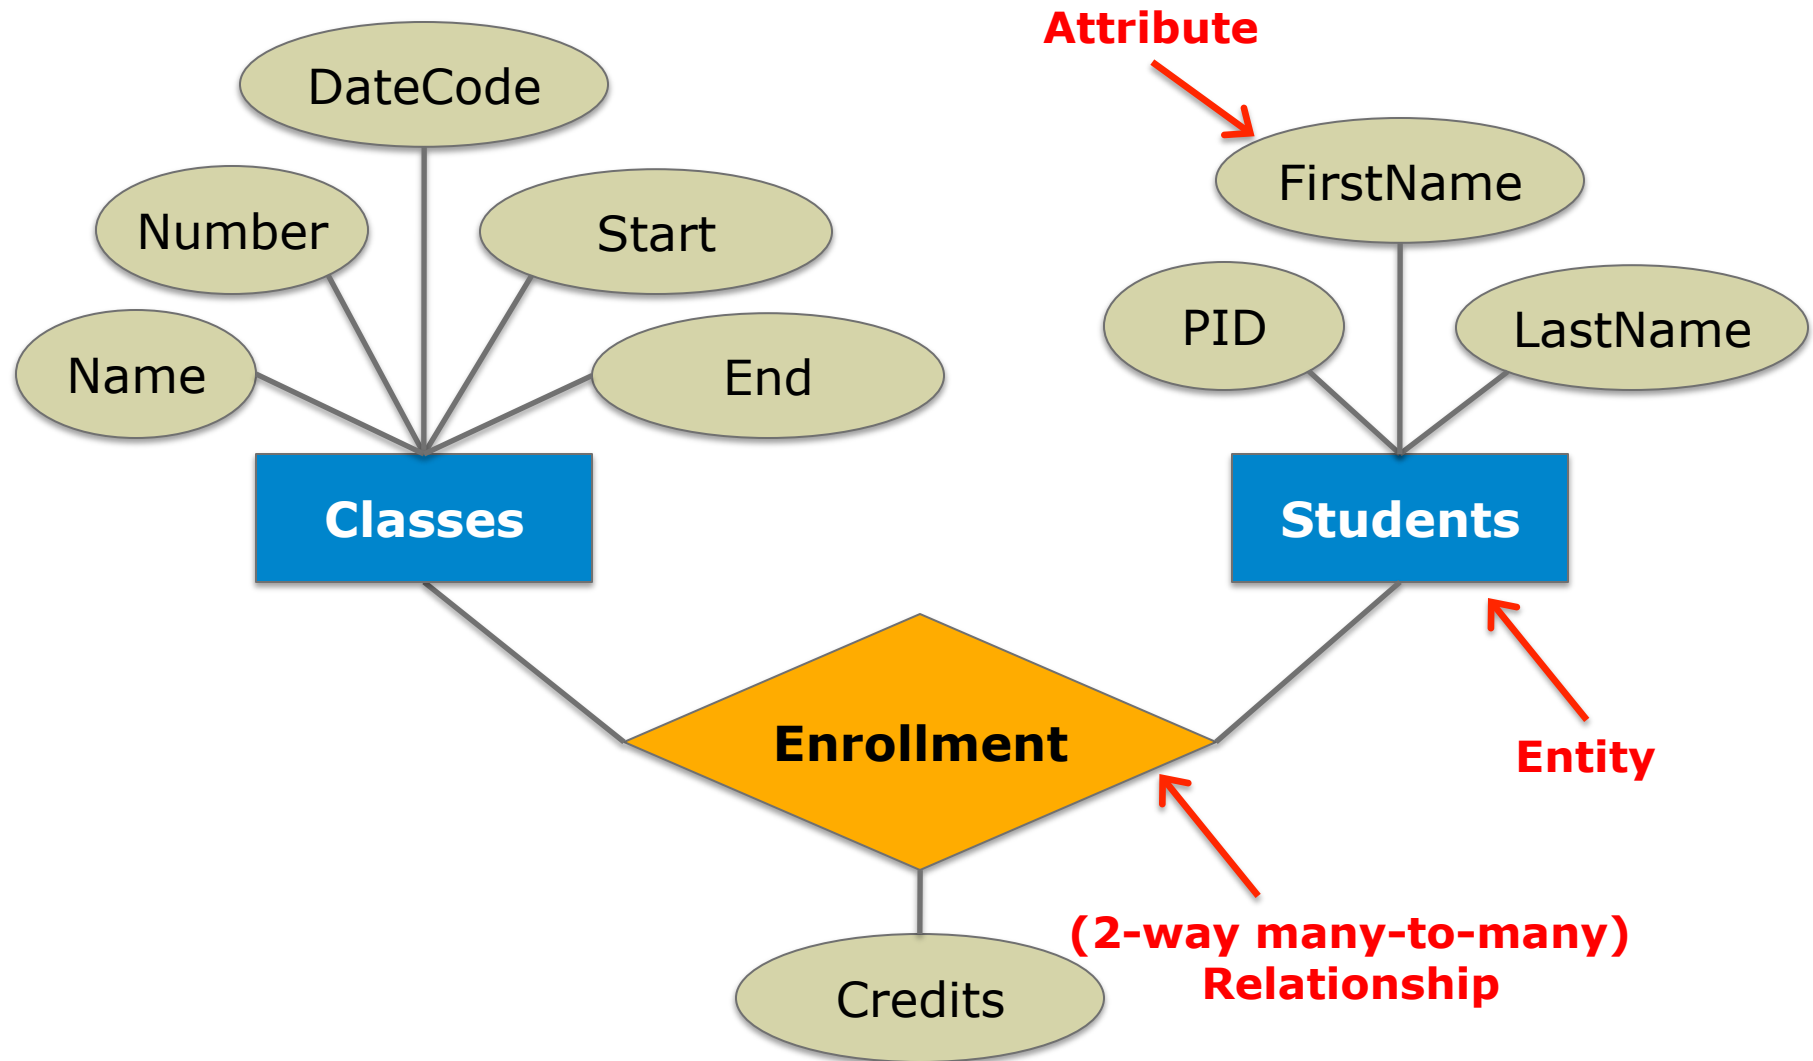
# Data Structure: Relational Model

## Example Problem:

- Represent the students classes of the CSE department in Winter, including the enrollment of students in classes.
- Students have pid, first name and last name.
- Classes have a name, a number, date code (TR, MW, MWF) and start/end time.
  - Dismiss the possibility of two Winter classes (or class sections) for the same course
- A student enrolls for a number of credits in a class.

## Solution:...

# Example 1a: E/R-Based Design



# E/R → Relational Schema:

## Basic Translation

- For every entity
  - create corresponding table
  - For each attribute of the entity, add a corresponding attribute in the table
  - Include an ID attribute in the table even if not in E/R
- For every many-to-many relationship
  - create corresponding table
  - For each attribute of the relationship, add a corresponding attribute in the table
  - For each referenced entity  $E_i$  include in the table a *required foreign key* attribute referencing ID of  $E_i$

# Sample relational database, per previous page's algorithm

**Classes**

id	name	number	date_code	start_time	end_time
1	Web stuff	MAS201	TuTh	2:00	3:20
2	Databases	CSE132A	TuTh	3:30	4:50
4	VLSI	CSE121	F	<i>null</i>	<i>null</i>

**Enrollment**

id	class	student	credits
1	1	1	4
2	1	2	3
3	4	3	4
4	1	3	3

**Students**

id	pid	first_name	last_name
1	8888888	John	Smith
2	1111111	Mary	Doe
3	2222222	<i>null</i>	Chen

# Declaration of schemas in SQL's Data Definition Language

```
CREATE TABLE classes (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    number      TEXT,  
    date_code   TEXT,  
    start_time  TIME,  
    end_time    TIME  
)
```

If we had "ID **INTEGER** PRIMARY KEY" we would be responsible for coming up with ID values. **SERIAL** leads to a counter that automatically provides ID values upon insertion of new tuples

Changed name from "end" to "end\_time" since "end" is reserved keyword

```
CREATE TABLE students (  
    ID          SERIAL PRIMARY KEY,  
    pid         INTEGER,  
    first_name  TEXT,  
    last_name   TEXT  
)
```

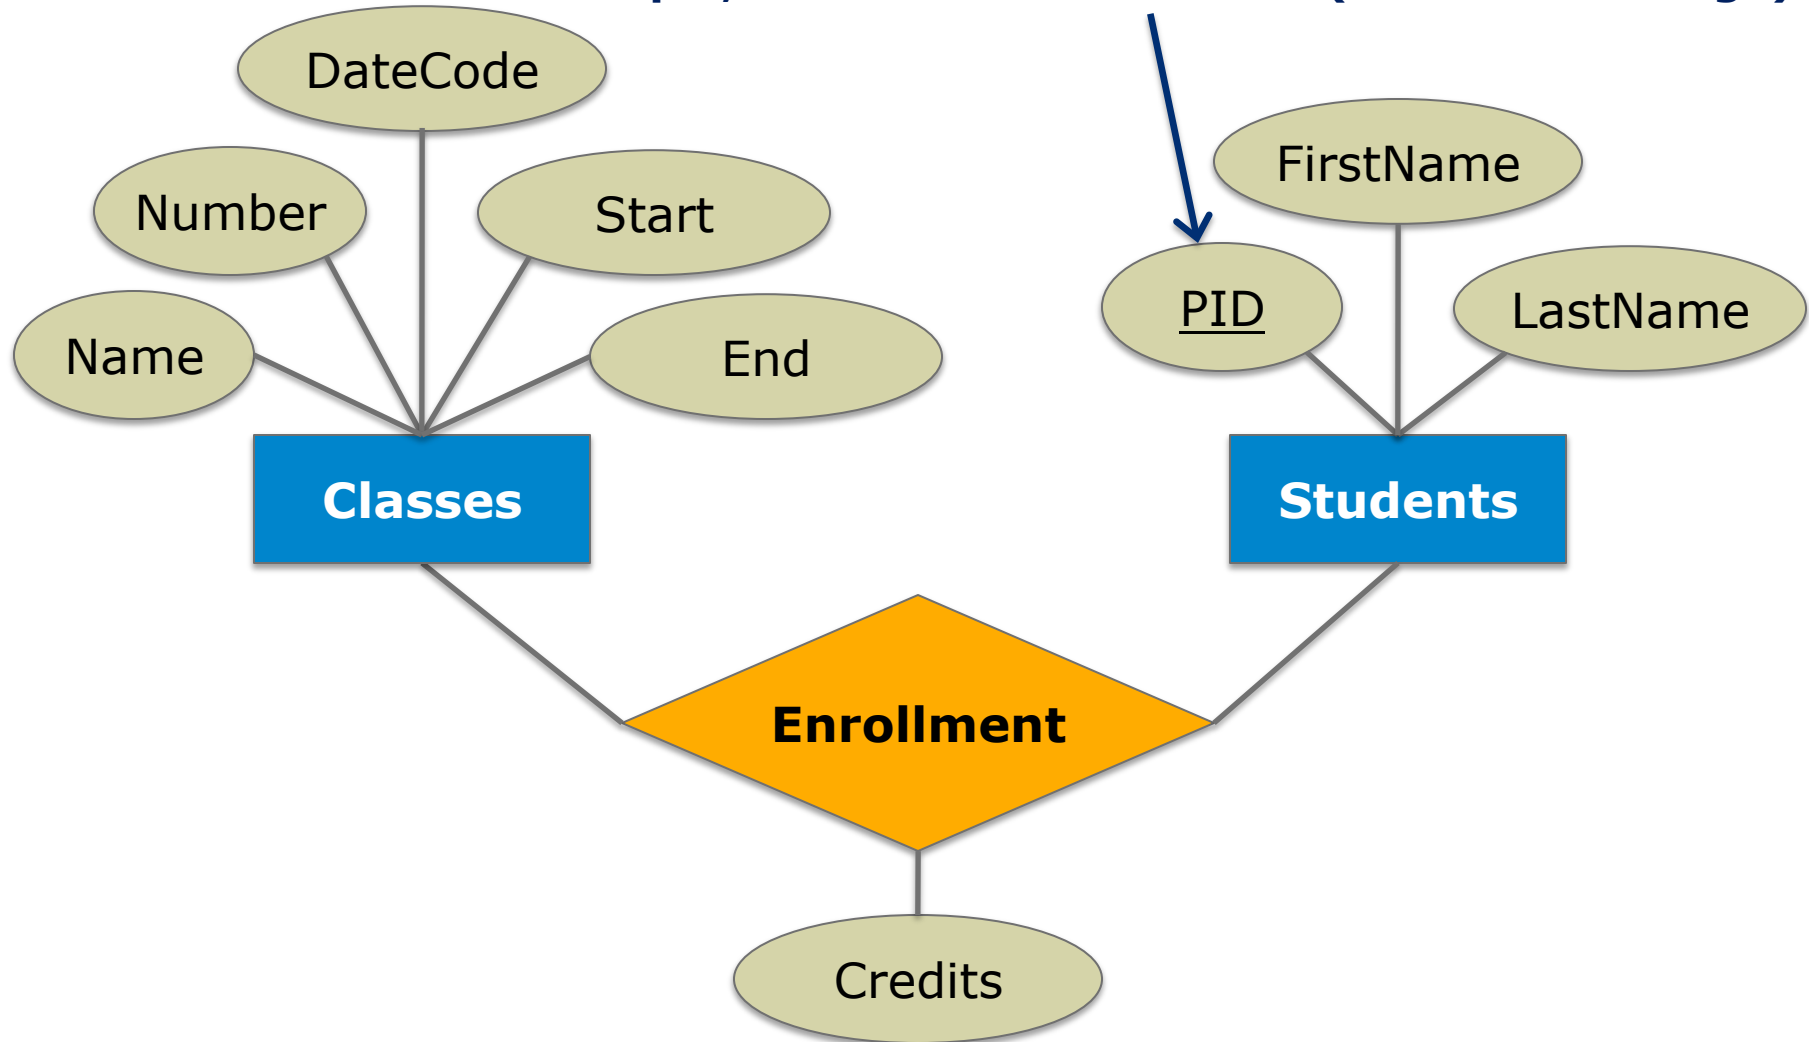
Foreign key declaration: Every value of **enrollment.class** must also appear as **classes.ID**

```
CREATE TABLE enrollment (  
    ID          SERIAL,  
    class       INTEGER REFERENCES classes (ID) NOT NULL,  
    student     INTEGER REFERENCES students (ID) NOT NULL,  
    credits     INTEGER  
)
```

Declaration of "required" constraint: **enrollment.student** cannot be null (notice, it would make no sense to have an enrollment tuple without a student involved)

# Example 1b: Using a semantic, immutable key

Assume that each PID (the id number on UCSD cards) is unique, not null and immutable (will never change)



# Example 1b: Sample, using the pid instead of the id to identify students

Classes

id	name	number	date_code	start_time	end_time
1	Web stuff	MAS201	TuTh	2:00	3:20
2	Databases	CSE132A	TuTh	3:30	4:50
4	VLSI	CSE121	F	<i>null</i>	<i>null</i>

Enrollment

id	class	student	credits
1	1	8888888	4
2	1	1111111	3
3	4	2222222	4
4	1	2222222	3

Students

<del>id</del>	pid	first_name	last_name
<del>1</del>	8888888	John	Smith
<del>2</del>	1111111	Mary	Doe
<del>3</del>	2222222	<i>null</i>	Chen

## Example 1b: Schema revisited, for using pid for students' primary key

```
CREATE TABLE classes (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    number      TEXT,  
    date_code   TEXT,  
    start_time  TIME,  
    end_time    TIME  
)  
  
CREATE TABLE students (  
    ID          SERIAL PRIMARY KEY,  
    pid         INTEGER PRIMARY KEY,  
    first_name  TEXT,  
    last_name   TEXT  
)  
  
CREATE TABLE enrollment (  
    ID          SERIAL,  
    class       INTEGER REFERENCES classes (ID) NOT NULL,  
    student     INTEGER REFERENCES students (pid) NOT NULL,  
    credits     INTEGER  
)
```



## ... some easy hands-on experience

- Install the Postgresql open source database
- For educational and management purposes use the pgAdmin client to define schemas, insert data,
- For managing and accessing the Postgresql server, use the pgAdmin graphical client
  - Right click on `Postgresql`, and select `Connect`
  - Right click on `Databases`, and select `New Database`
  - Enter a new name for the database, and click `Okay`
  - Highlight the database, and select `Tools -> Query Tool`
  - Write SQL code (or open the examples), and select `Query -> Execute`

# Creating a schema and inserting some data

- Open file enrollment.sql
- Copy and paste its CREATE TABLE and INSERT commands in the Query Tool
- Run it – you now have the sample database!
- Run the first 3 SELECT commands to see the data you have in the database
  - You can run a command by highlighting it with the cursor and click run

## Example 2a

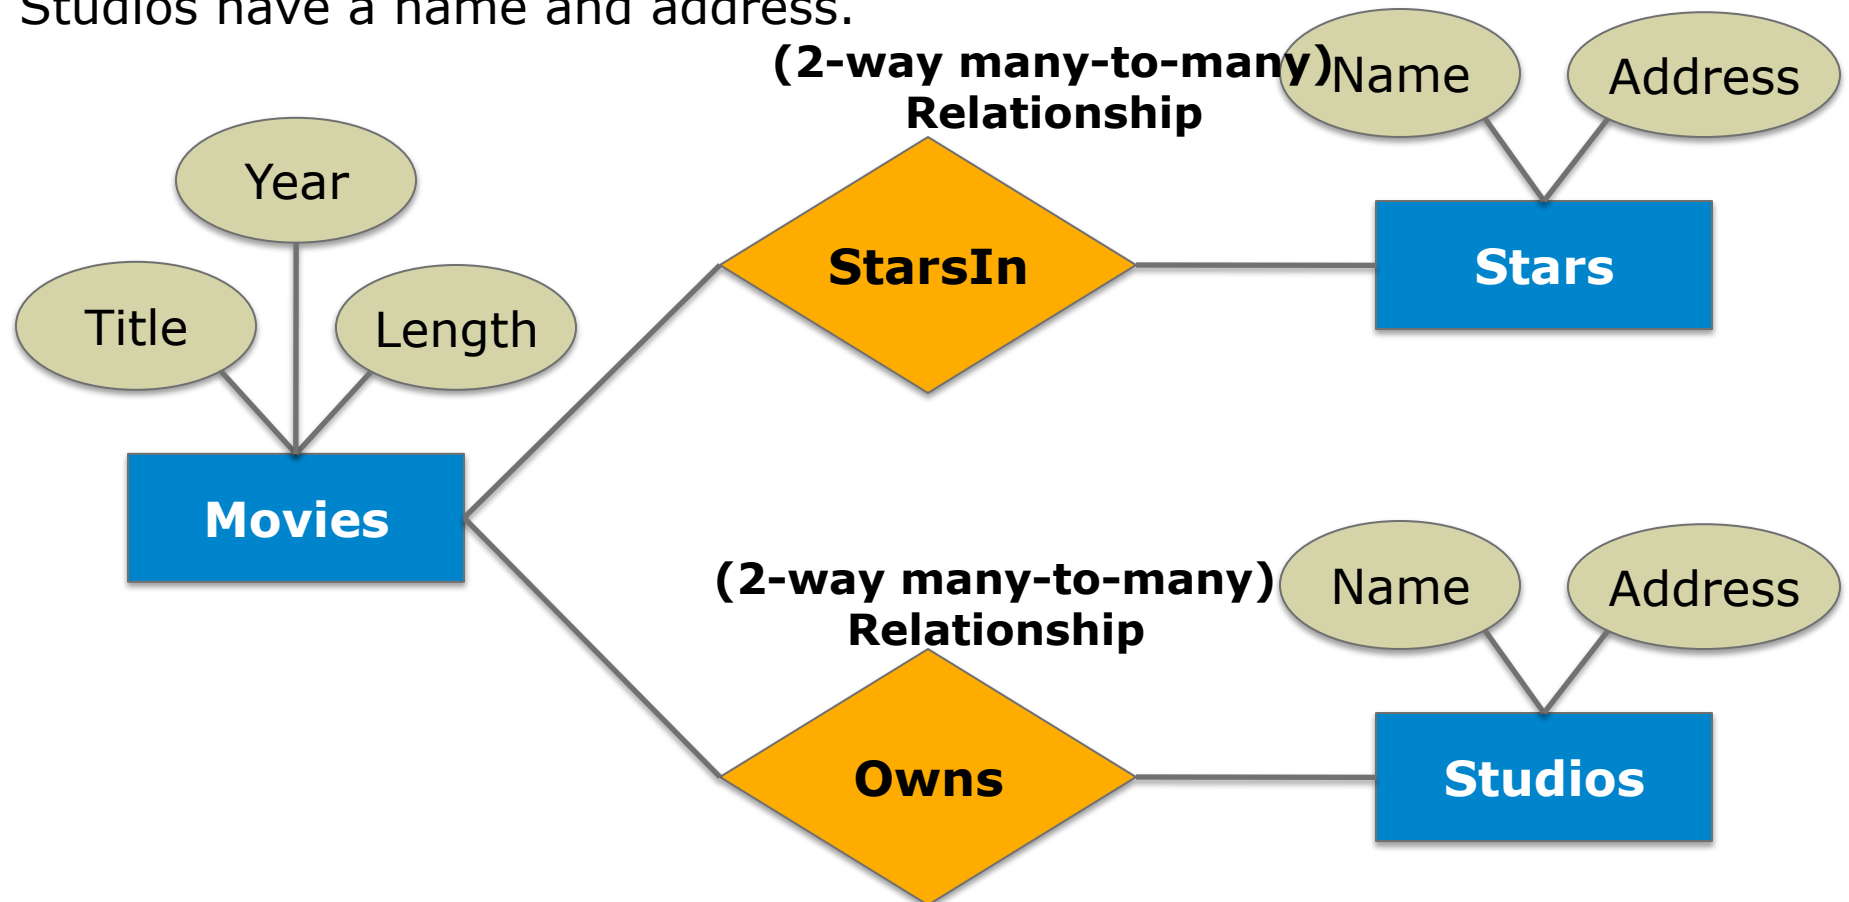
Movies have a title, a year of release and length (in minutes).

Actors have names and address.

Actors appear in movies.

A movie is (co-)owned by studios.

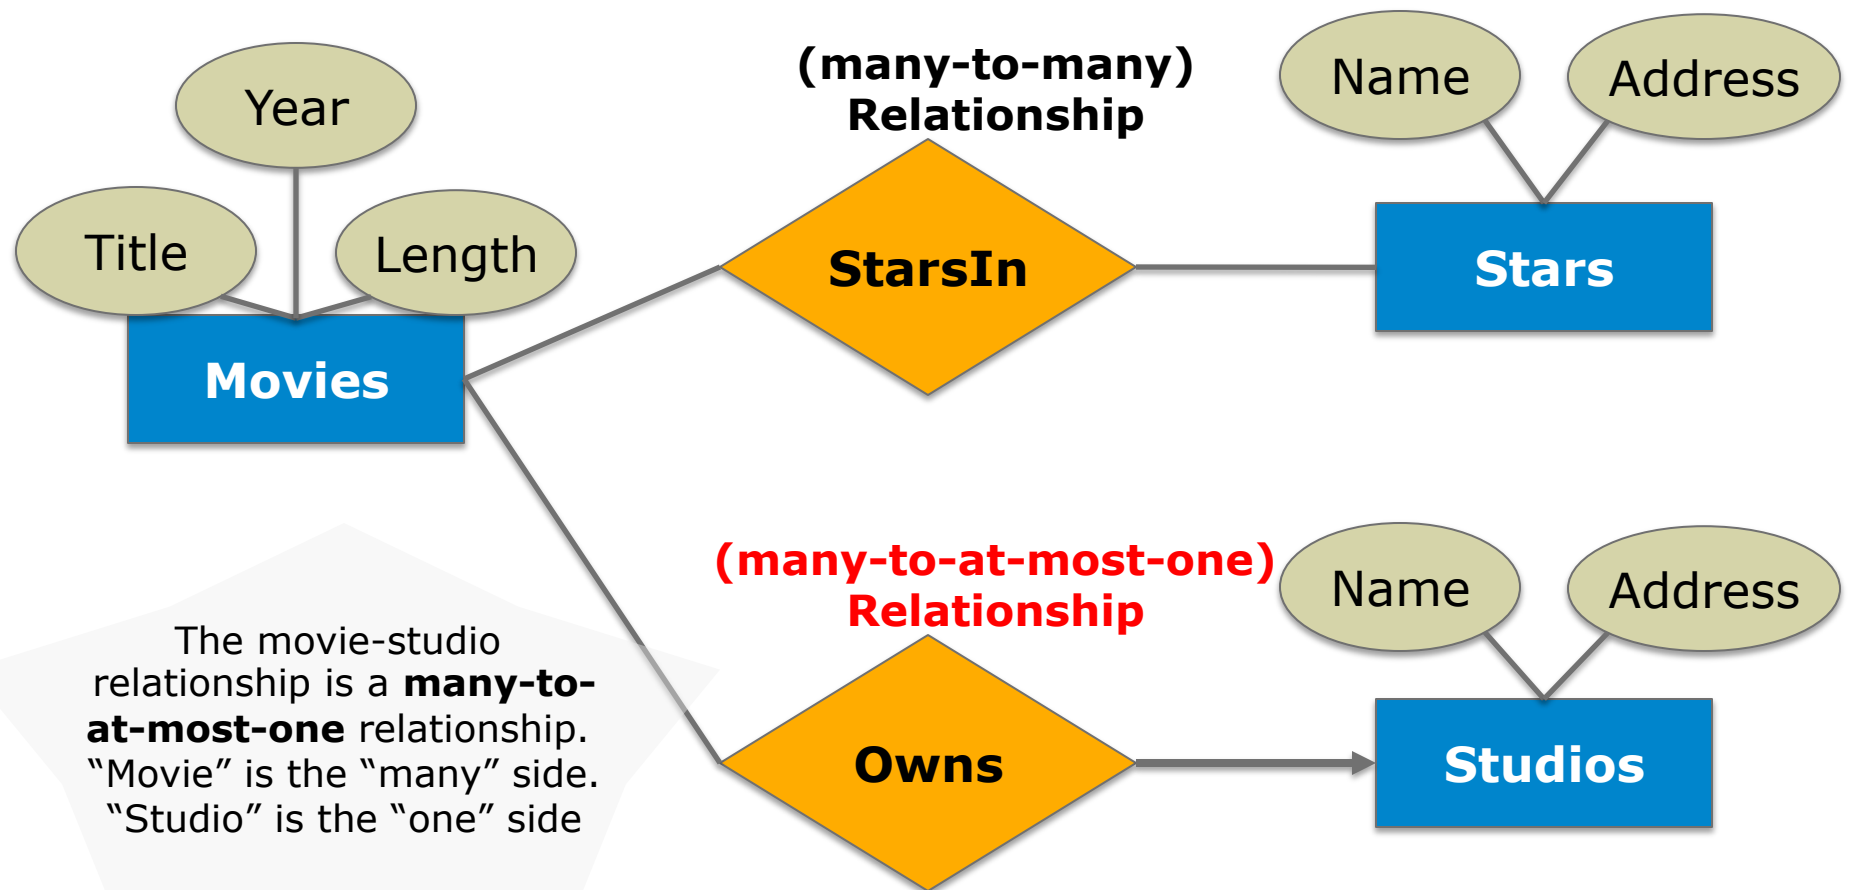
Studios have a name and address.



```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    title       TEXT,  
    year        INTEGER,  
    length      INTEGER,  
)  
CREATE TABLE stars (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)  
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)  
CREATE TABLE starsin (  
    ID          SERIAL,  
    movie       INTEGER REFERENCES movies (ID) NOT NULL,  
    star        INTEGER REFERENCES stars (ID) NOT NULL  
)  
CREATE TABLE ownership (  
    ID          SERIAL,  
    movie       INTEGER REFERENCES movies (ID) NOT NULL,  
    owner       INTEGER REFERENCES studios (ID) NOT NULL  
)
```

# Example 2b: many-to-at-most-one relationship

Modification to Example 2a:  
A movie is owned by **at most one** studio.



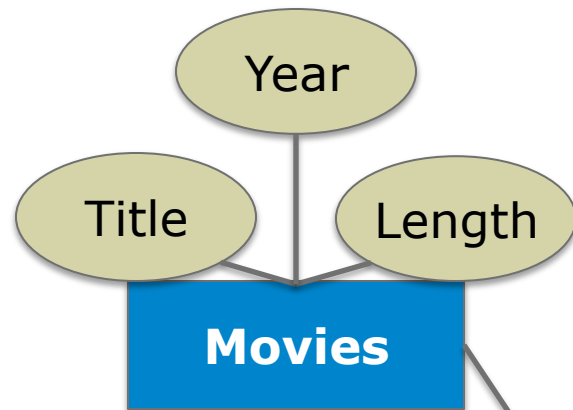
## E/R → Relational: Translation revisited for many-to-at-most-one relationship

- For every entity, do the usual...
- For every **many-to-many** relationship, do the usual...
- For every **2-way many-to-at-most-one** relationship, where
  - $E_m$  is the “many” side
  - $E_o$  is the “one” side (pointed by the arrow)
  - **do not** create table, instead:
  - In the table corresponding to  $E_m$  add a (non-required, i.e., potentially NULL) foreign key attribute referencing the ID of the table corresponding to  $E_o$

```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    title       TEXT,  
    year        INTEGER,  
    length      INTEGER,  
    owner       INTEGER REFERENCES studios (ID)  
)  
CREATE TABLE stars (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)  
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)  
CREATE TABLE starsin (  
    ID          SERIAL,  
    movie       INTEGER REFERENCES movies (ID) NOT NULL,  
    star        INTEGER REFERENCES stars (ID) NOT NULL  
)
```



# Example 2c: many-to-exactly-one relationship

Modification to Example 2a:  
A movie **must** be owned by **one** studio.



```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    title       TEXT,  
    year        INTEGER,  
    length      INTEGER,  
    owner       INTEGER REFERENCES  
                studios (ID) NOT NULL  
)
```

**(many-to-exactly-one)  
Relationship**




The movie-studio relationship is a **many-to-exactly-one** relationship. "Movie" is the "many" side. "Studio" is the "one" side




## Example 2d: one-to-one relationship

Consider a database of heterosexual couples  
(we neglect homosexual couples, amazons and Warren Jeffs followers)



```
CREATE TABLE couple (  
    husband INTEGER REFERENCES  
        females (ID) NOT NULL UNIQUE,  
    wife INTEGER REFERENCES  
        males (ID) NOT NULL UNIQUE  
)
```

## Example 3: 3-Way Many-to-Many Relationship




- A studio has contracted with a particular star to act in a particular movie
  - No ownership of movies by studios

---

```
CREATE TABLE contract (  
    ID                SERIAL,  
    movie             INTEGER REFERENCES movies (ID) NOT NULL,  
    star              INTEGER REFERENCES stars (ID) NOT NULL,  
    owner             INTEGER REFERENCES studios (ID) NOT NULL,  
    fee               INTEGER  
)
```

# Example 4a : Self-Relationships with Roles

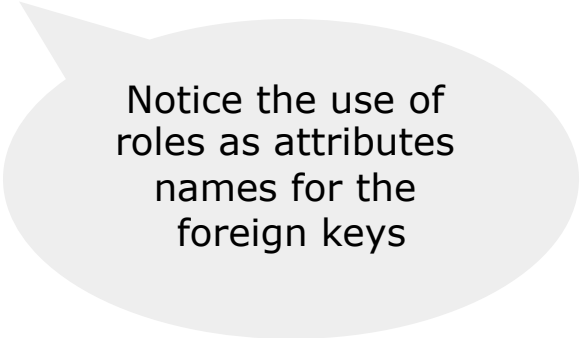
Twitter Use Case



---

```
CREATE TABLE persons (  
    ID          SERIAL PRIMARY KEY,  
    ...  
)
```


```
CREATE TABLE following (  
    ID          SERIAL,  
    follows     INTEGER REFERENCES persons (ID) NOT NULL,  
    isFollowed  INTEGER REFERENCES persons (ID) NOT NULL  
)
```



Notice the use of  
roles as attributes  
names for the  
foreign keys

# Example 4b : Self-Relationships with Roles

Prequels and Sequels



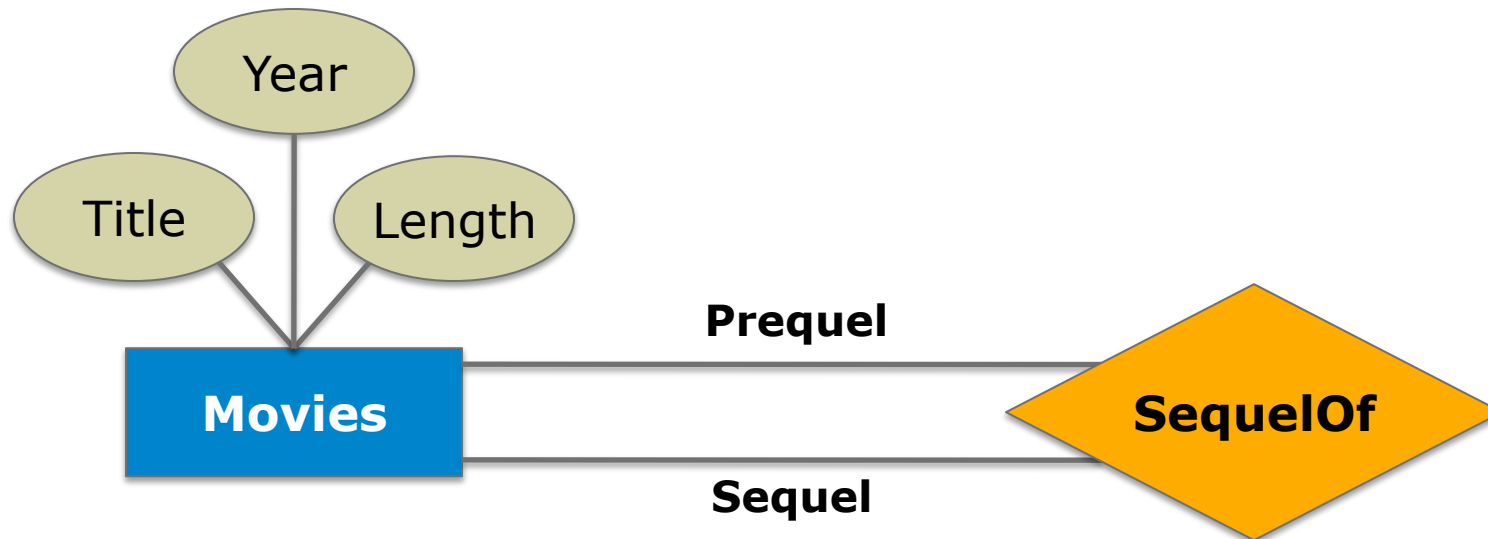
---

```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    ...  
)
```

```
CREATE TABLE sequelof (  
    ID          SERIAL,  
    prequel     INTEGER REFERENCES movies (ID) NOT NULL,  
    sequel      INTEGER REFERENCES movies (ID) NOT NULL  
)
```

## Example 4b : Self-Relationships with Roles – Questions on Meaning

What exactly are the prequel-sequel pairs?



“Terminator II: Judgment Day” is a sequel of “Terminator”

“Terminator III: Raise of the Machines” is a sequel of  
“Terminator II: Judgment Day”

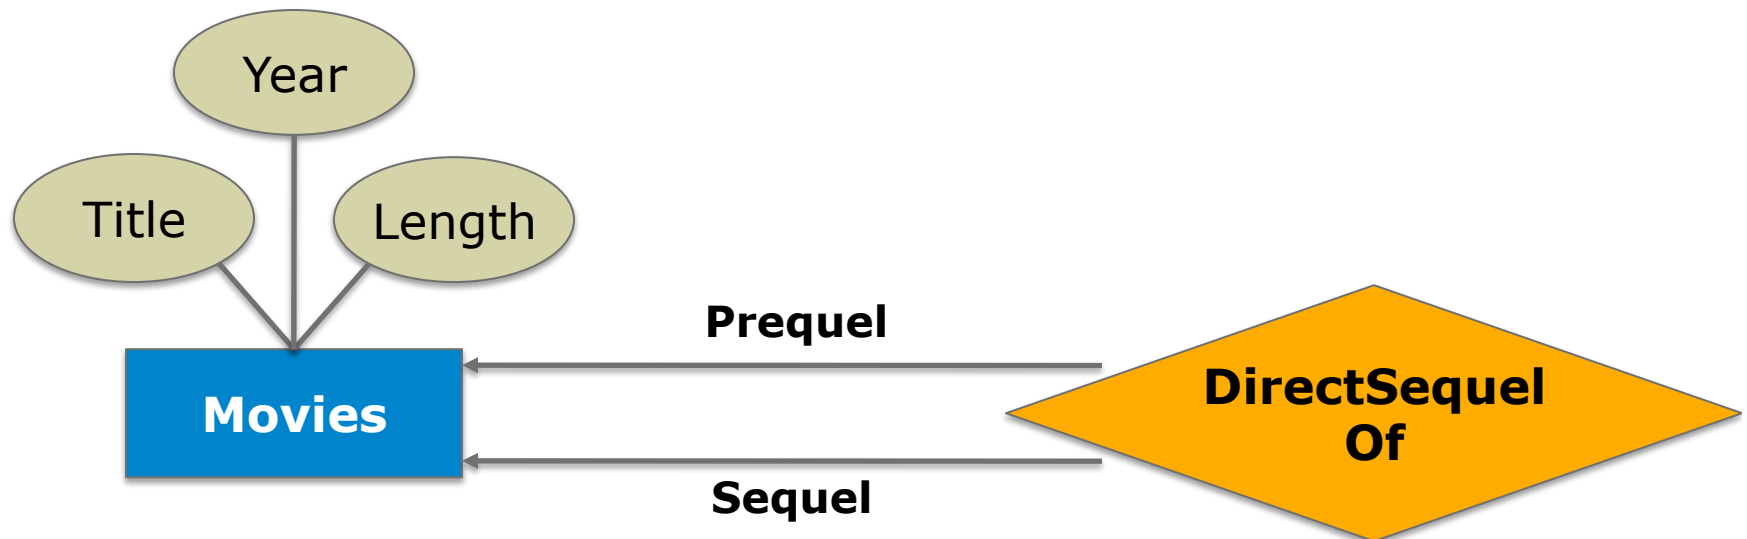
**Is “Terminator III: Raise of the Machines” a sequel of  
“Terminator” ?**



## Example 4c : Interpreting sequels non-transitively

Is “Terminator III: Raise of the Machines”  
a **direct** sequel of “Terminator” ? **NO**

A movie has at most one direct “prequel” and at most one direct “sequel”



Modeling movie sequels by “DirectSequelOf” is preferable in OLTP to using transitive “SequelOf”

- A lesson about good (OLTP?) database design:
- Good designs avoid redundancy.
  - No stored piece of data should be inferable from other stored pieces of data

# To be Redundant or Not to be?

## NOT


- Too many Friends-of-Friends
  - Even more Friends-of-Friends-of-Friends
    - If “Six Degrees of Separation” is true, the 6-step friends is not even saying anything
- A database with derivative data is harder to maintain

## YES


- Some derivations, interesting to OLAP, are too expensive to compute live
- If OLAP, maintenance is not primary concern

# Self-relationships without roles


Twitter “followship” is a self-relationship with roles



Facebook “friendship” is a self-relationship without real roles



# A case where redundancy may be welcome



```
CREATE TABLE friend (  
    subject INTEGER REFERENCES user (ID) NOT NULL,  
    object INTEGER REFERENCES user (ID) NOT NULL  
)
```

If Subject is Facebook friend of Object,  
then Object is Facebook friend of Subject.

Is it redundant to explicitly represent both facts in "friend"?

Yes, but makes some queries much easier and faster.

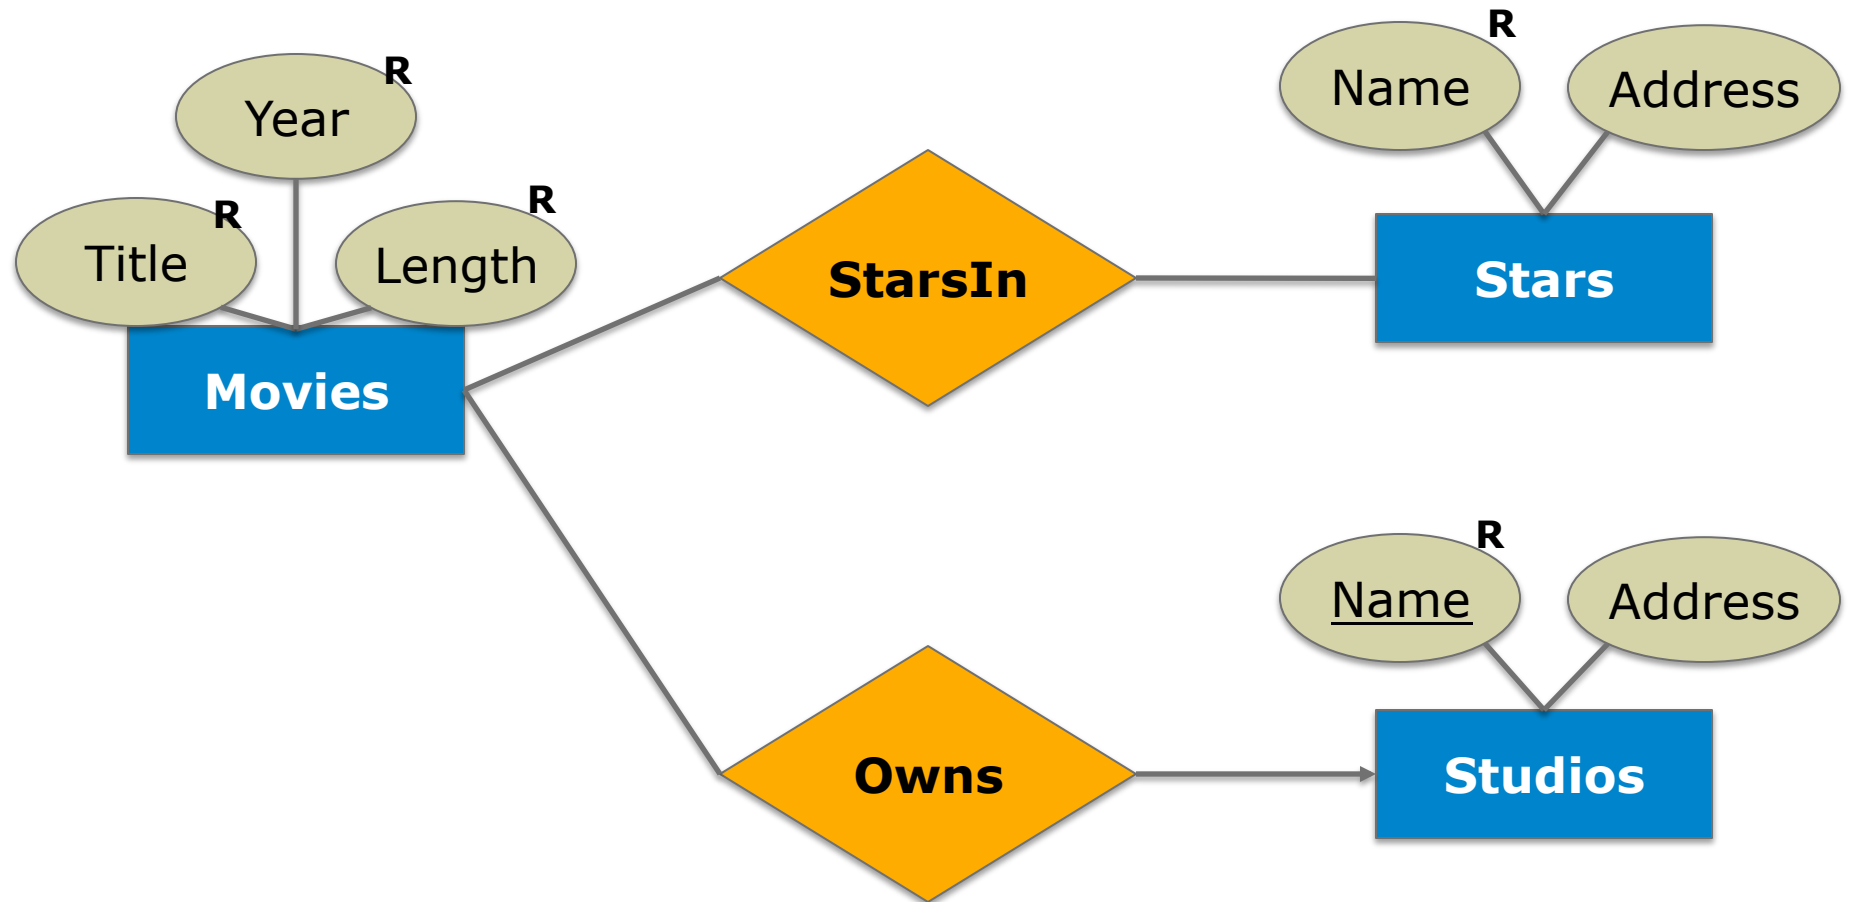
# **Designing Schemas Using Entity-Relationship modeling**

Additional Topics

## Example 5a: Constraints: uniqueness; required attributes

In addition to Example 2b's assumptions, let us also assume that:

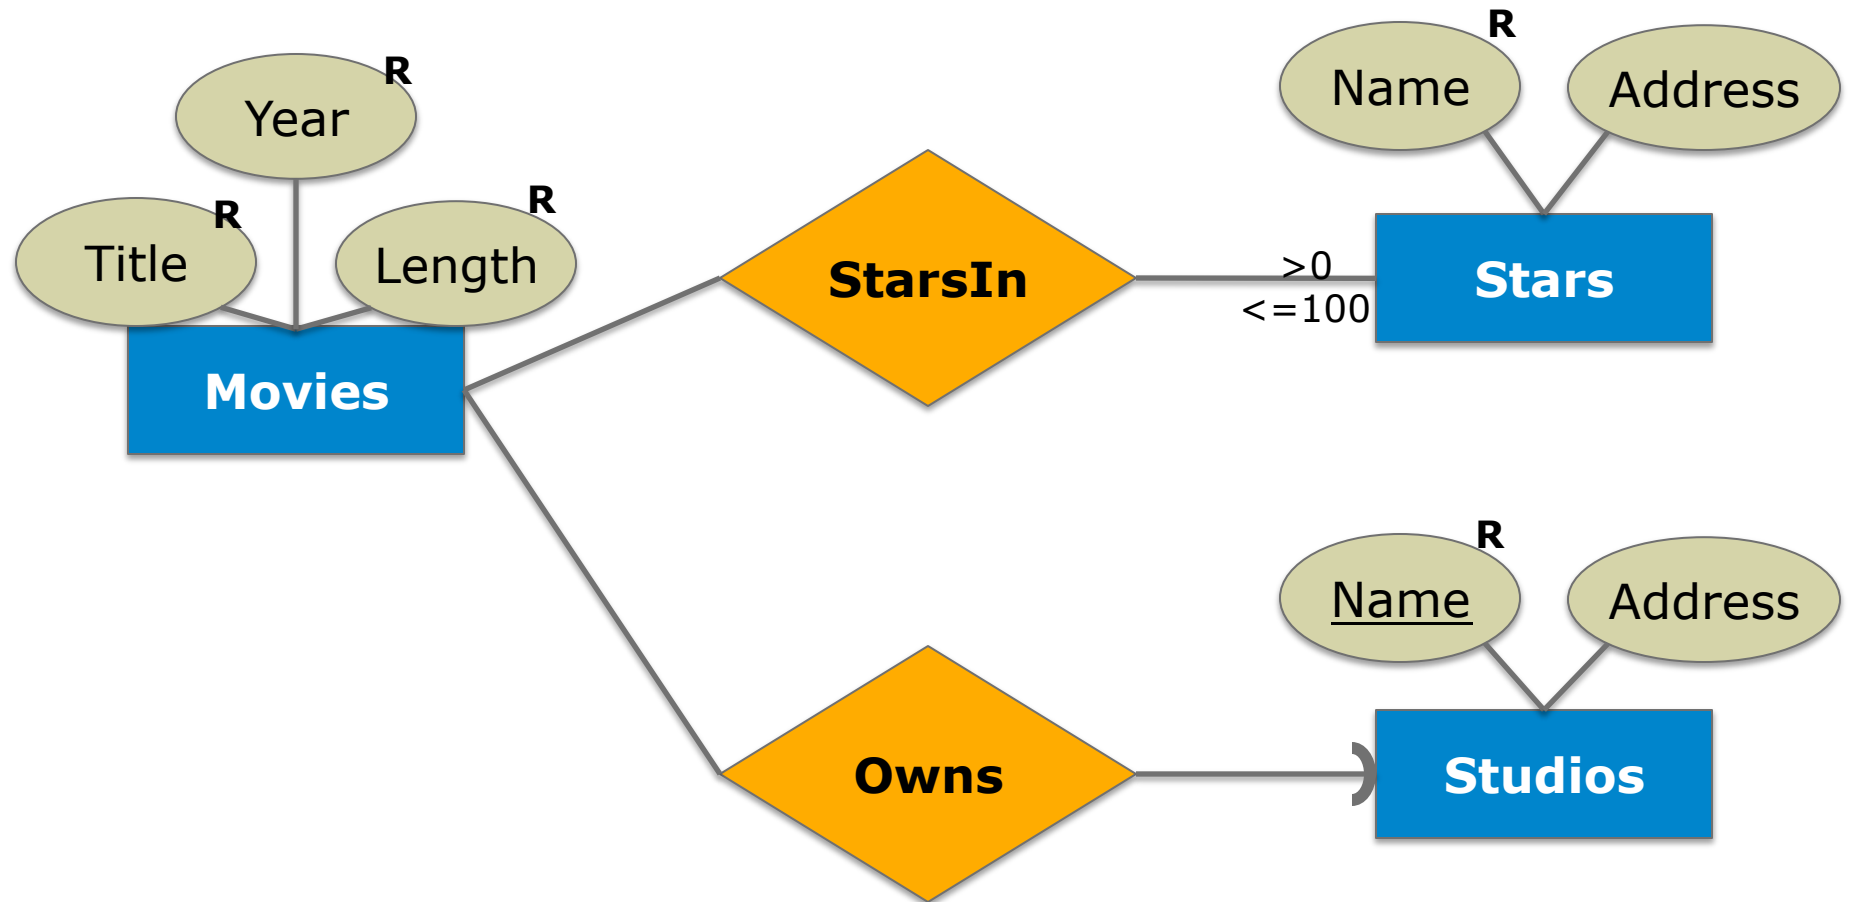
- title, year, length, star name and studio name are required attributes of the respective entities
  - default is that an attribute value may be **null**
- studios have unique names, i.e., no two studios may have the same name



## Example 5b: Constraints: Required relationship; cardinality ranges

In addition to Example 2c's assumptions, let us also assume that:

- a movie is owned by **exactly one** studio
  - so far we had not assumed that the owning studio has to be known (not null)
- a movie must have at least one actor and no more than 100



# SQL Schema for Examples 5a, 5b

```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    title       TEXT NOT NULL,  
    year        INTEGER NOT NULL,  
    length      INTEGER NOT NULL,  
    owner       INTEGER REFERENCES studios (ID) NOT NULL  
)  
  
CREATE TABLE stars (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT NOT NULL,  
    address     TEXT  
)  
  
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT NOT NULL UNIQUE,  
    address     TEXT  
)  
  
CREATE TABLE starsin (  
    ID          SERIAL,  
    movie       INTEGER REFERENCES movies (ID) NOT NULL,  
    star        INTEGER REFERENCES stars (ID) NOT NULL  
)
```



# A sample database

**stars**

ID	name	address
1	Al Pacino	New York, NY
2	Harrison Ford	Beverly Hills, CA
3	Tom Hanks	Santa Monica, CA

**studios**

ID	name	address
1	20th Century Fox	Century City, CA
2	Paramount Productions	Hollywood, CA
3	Universal Pictures	Universal City, CA

**starsin**

ID	movie	star
1	1	3
2	2	1
3	3	2
4	4	1

**movies**

ID	title	year	length	owner
1	Forrest Gump	1994	142	2
2	The Godfather	1972	175	2
3	Star Wars	1977	121	1
4	Scent of a Woman	1992	157	3

# Why do we want constraints? What happens when they are violated?

- Protect the database from erroneous data entry
- Prevent database states that are inconsistent with the rules of the business process you want to capture
- Whenever you attempt to change (insert, delete, update) the database in a way that violates a constraint the database will prevent the change
  - Try it out on the sample databases of the class page

# Some constraints are not implemented by some SQL database systems

- Consider the cardinality constraint that a movie has between 1 and 100 actors.
- The SQL standard provides a way, named CHECK constraints, to declare such
  - its specifics will make more sense once we have seen SQL queries
- However, no open source database implements the CHECK constraints.
- Project Phase I: Introduce such constraints on your E/R, despite the fact that you will not be able to translate them to the SQL schema

# Vice versa: SQL allows some constraints that are not in plain E/R

Notable cases:

- Attribute value ranges
  - Example: Declare that the year of movies is after 1900
- Multi-attribute UNIQUE
  - Example: Declare that the (title, year) attribute value combination is unique

# Added constraints of previous slide

```
CREATE TABLE movies (  
    ID          SERIAL PRIMARY KEY,  
    title       TEXT NOT NULL,  
    year        INTEGER NOT NULL CHECK (year > 1900),  
    length      INTEGER NOT NULL,  
    owner       INTEGER REFERENCES studios (ID) NOT NULL,  
    UNIQUE (title, year)  
)  
  
CREATE TABLE stars (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT NOT NULL,  
    address     TEXT  
)  
  
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT NOT NULL UNIQUE,  
    address     TEXT  
)  
  
CREATE TABLE starsin (  
    ID          SERIAL,  
    movie       INTEGER REFERENCES movies (ID) NOT NULL,  
    star        INTEGER REFERENCES stars (ID) NOT NULL  
)
```

## Example 6: tricky flavors of one-to-one relationships

Assume that a president manages exactly one studio and a studio may have at most one president.

Notice: a studio may not have a president but in order to be a president one has to manage a studio.



# 1<sup>st</sup> candidate

```
CREATE TABLE presidents (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    age         INTEGER,  
)
```

```
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)
```

```
CREATE TABLE management (  
    manager INTEGER REFERENCES presidents (ID) NOT NULL UNIQUE  
    manages INTEGER REFERENCES studios (ID) NOT NULL UNIQUE  
)
```

One may be a president  
WITHOUT managing any  
studio  
=> This design fails to  
capture a given constraint

## 2<sup>nd</sup> candidate

```
CREATE TABLE presidents (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    age         INTEGER,  
    manages     INTEGER REFERENCES studios (ID) NOT NULL UNIQUE  
)
```

```
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT  
)
```

Guarantees that  
in order to be  
president, one  
has to manage a  
studio

Guarantees that  
no two presidents  
may manage the  
same studio

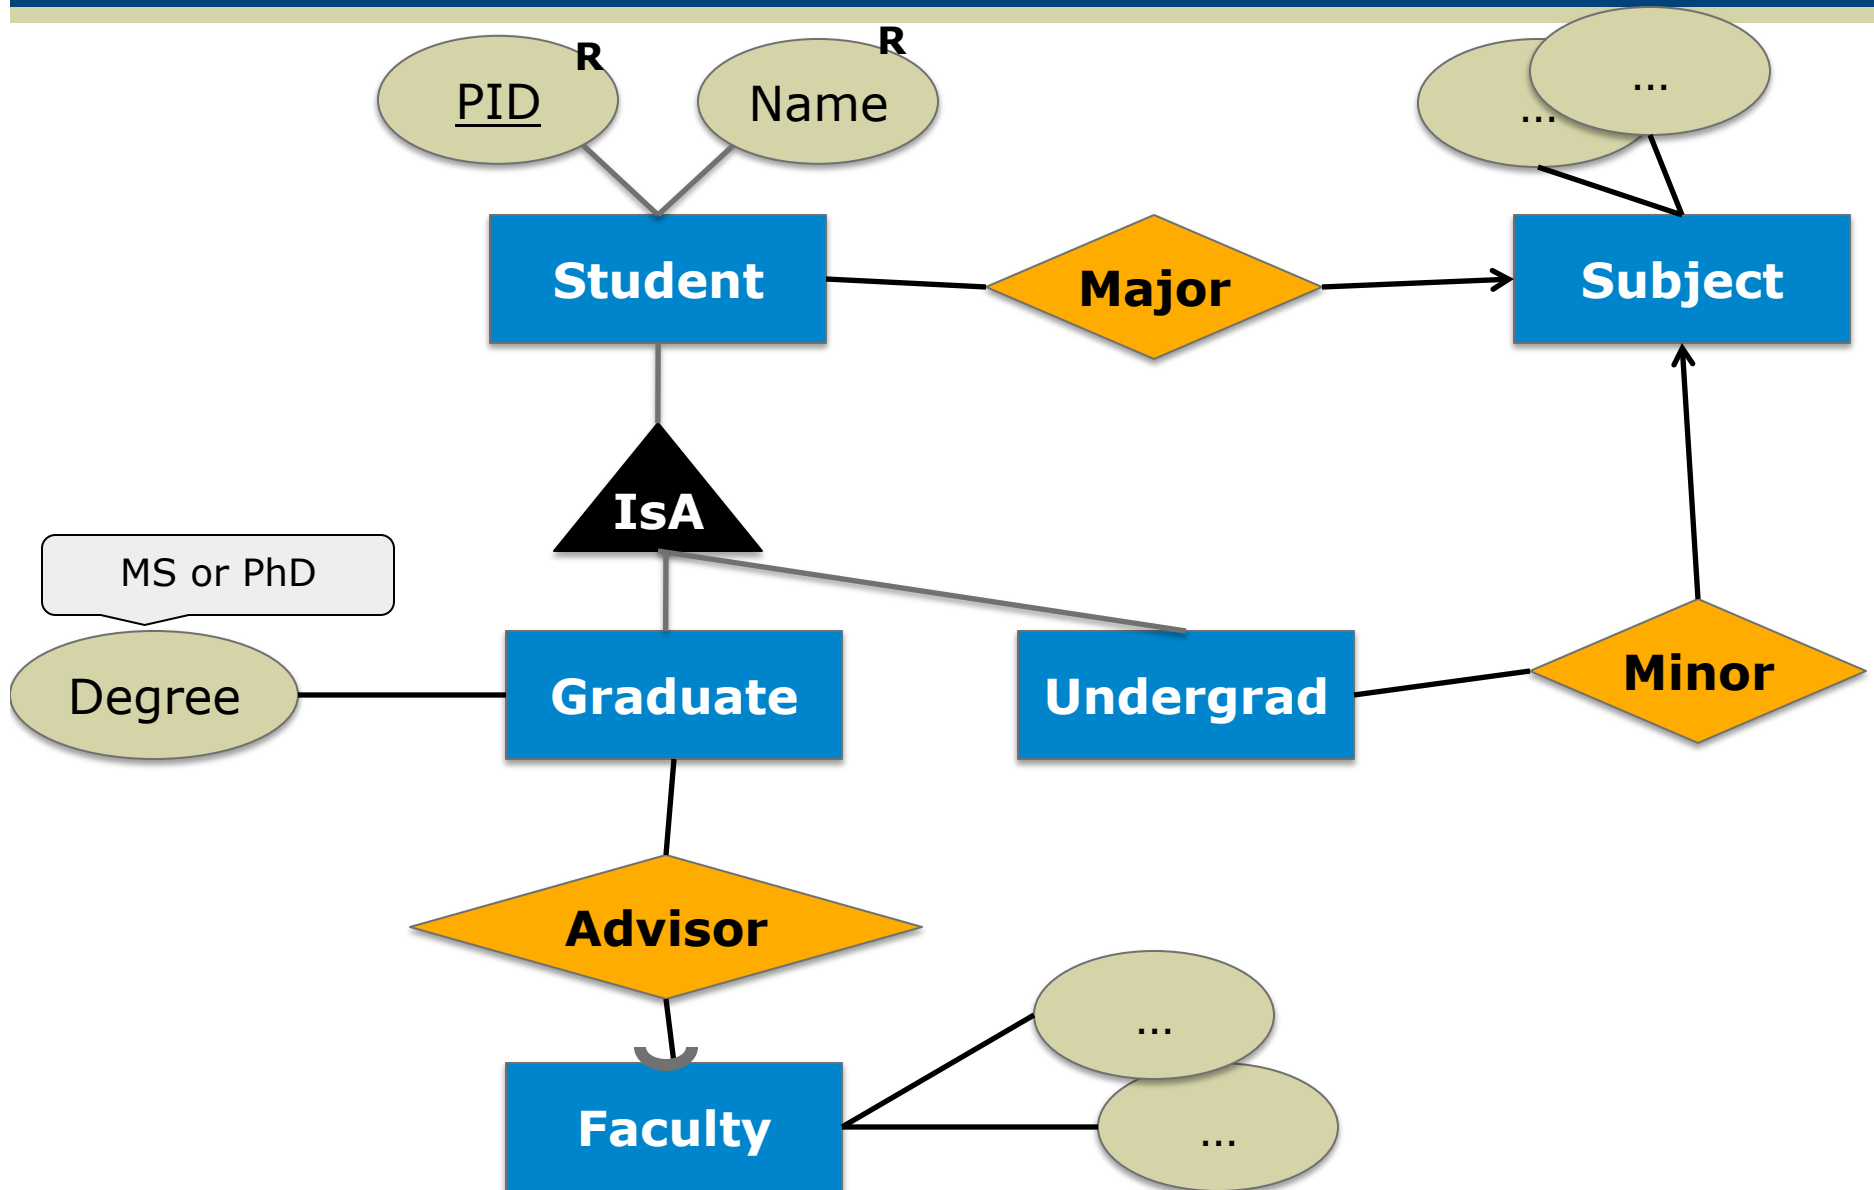


# 3rd candidate

3rd candidate is also not preferred. Why? What constraint it misses?

```
CREATE TABLE presidents (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    age         INTEGER  
)  
  
CREATE TABLE studios (  
    ID          SERIAL PRIMARY KEY,  
    name        TEXT,  
    address     TEXT,  
    managedBy   INTEGER REFERENCES presidents (ID) UNIQUE  
)
```

# Example 6: Subclassing



# Schemas for subclassing:

## Candidate 1

```
CREATE TABLE student(  
    ID        SERIAL PRIMARY KEY,  
    pid       TEXT NOT NULL UNIQUE,  
    name      TEXT NOT NULL,  
    major     INTEGER REFERENCES subject(ID)  
)  
CREATE TABLE undergrad(  
    studentid  INTEGER REFERENCES student(ID) NOT NULL,  
    minor      INTEGER REFERENCES subject(ID)  
)  
CREATE TABLE graduate(  
    studentid  INTEGER REFERENCES student(ID) NOT NULL,  
    degree     TEXT NOT NULL CHECK (degree="PhD" OR degree="MS"),  
    advisor    INTEGER REFERENCES faculty(ID) NOT NULL  
)  
CREATE TABLE subject(  
    ID        SERIAL PRIMARY KEY,  
    ...  
)  
CREATE TABLE faculty(  
    ID        SERIAL PRIMARY KEY,  
    ...  
)
```

+ captures constraints  
- Information about graduates is spread on two tables  
- Creating a report about students is a tricky query  
To appreciate the above wait till we discuss SQL

# Schemas for subclassing: Candidate 2

```
CREATE TABLE student(  
    ID          SERIAL PRIMARY KEY,  
    pid         TEXT NOT NULL UNIQUE,  
    name        TEXT NOT NULL,  
    kind        CHAR(1) CHECK (kind='U' OR kind='S'),  
    major       INTEGER REFERENCES subject(ID),  
    minor       INTEGER REFERENCES subject(ID),  
    degree      TEXT CHECK (degree="PhD" OR degree="MS"),  
    advisor     INTEGER REFERENCES faculty(ID)  
)  
CREATE TABLE subject(  
    ID          SERIAL PRIMARY KEY,  
    ...  
)  
CREATE TABLE faculty(  
    ID          SERIAL PRIMARY KEY,  
    ...  
)
```

-misses constraints  
E.g., notice that it does not capture that a graduate student must have an advisor since we had to make the advisor attribute non-required in order to accommodate having undergraduates in the same table

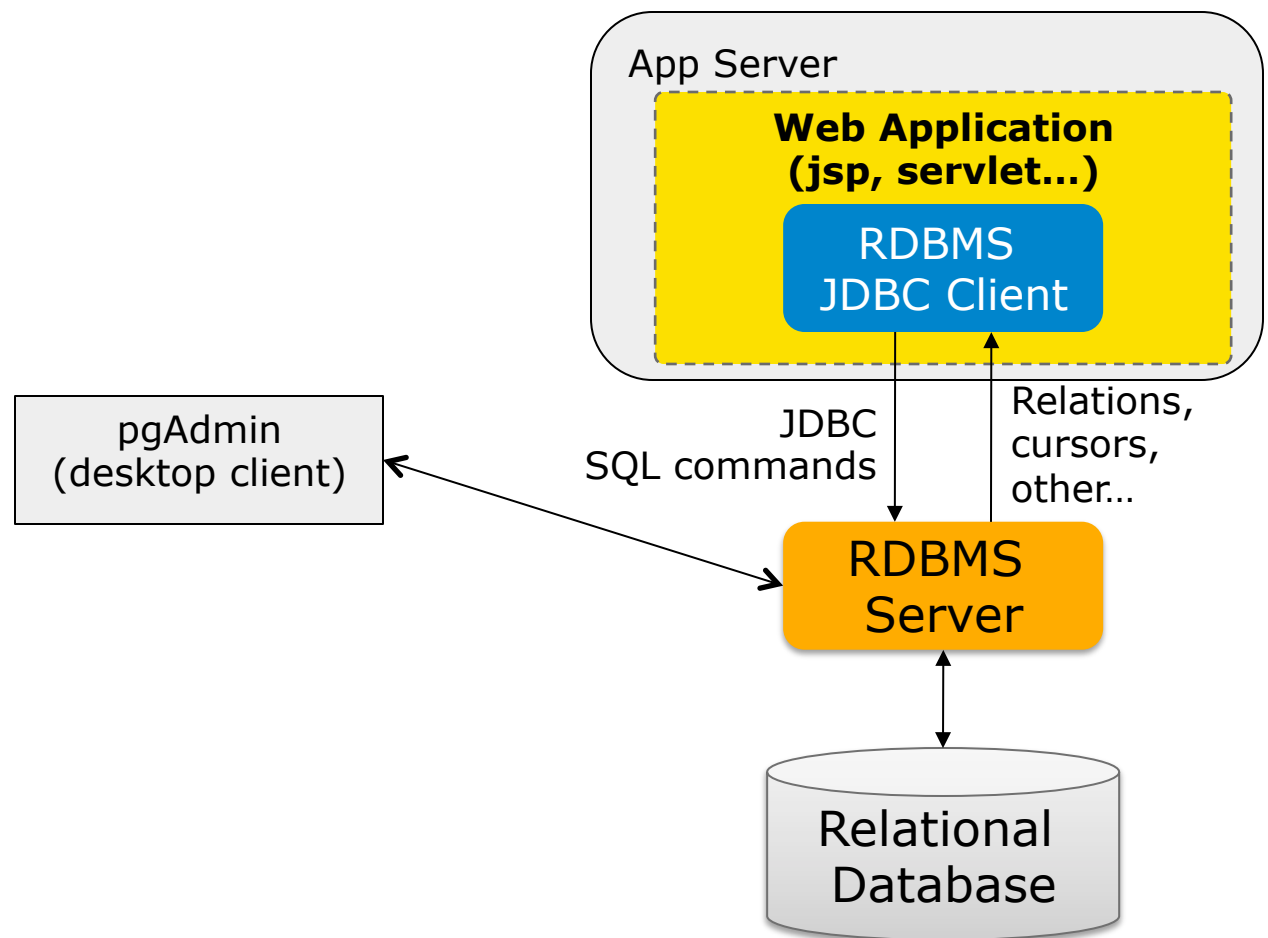
# Not covered E/R features

- Weak entities
  - double-lined entities and relationships
- Many-to-Many-to-One 3-way (or more) relationships
- Necessary participation of entity in relationship
- ... more

# **Programming on Databases with SQL**

# Writing programs on databases: JDBC

- How client opens connection with a server
- How access & modification commands are issued
- ...



## ... some easy hands-on experience

- Install the Postgresql open source database
- For educational and management purposes use the pgAdmin client to define schemas, insert data,
  - See online instructions
- For managing and accessing the Postgresql server, use the pgAdmin graphical client
  - Right click on `Postgresql`, and select `Connect`
  - Right click on `Databases`, and select `New Database`
  - Enter a new name for the database, and click `Okay`
  - Highlight the database, and select `Tools -> Query Tool`
  - Write SQL code (or open the examples), and select `Query -> Execute`



# Creating a schema and inserting some data

- Open file enrollment.sql
- Copy and paste its CREATE TABLE and INSERT commands in the Query Tool
- Run it – you now have the sample database!
- Run the first 3 SELECT commands to see the data you have in the database

# Access (Query) & Modification Language: SQL

- SQL
  - used by the database user
  - **declarative**: we only describe **what** we want to retrieve
  - based on tuple relational calculus
- The result of a query is a table
- Internal Equivalent of SQL: Relational Algebra
  - used internally by the database system
  - **procedural** (operational): describes **how** query is executed
- The solutions to the following examples are on the class page download

# SQL: Basic, single-table queries

- Basic form

```
SELECT   $r.A_1, \dots, r.A_N$   
FROM R r
```

```
WHERE <condition>
```

- **WHERE** clause is optional
- Have tuple variable  $r$  range over the tuples of  $R$ , qualify the ones that satisfy the (boolean) condition and return the attributes  $A_1, \dots, A_N$

Find first names and last names of all students

```
SELECT s.first_name, s.last_name  
FROM  students s;
```

---

Display all columns of all students whose first name is John; project all attributes

```
SELECT s.id, s.pid, s.first_name,  
        s.last_name  
FROM students s  
WHERE s.first_name = 'John'
```

... and its shorthand form

```
SELECT *  
FROM  students s  
WHERE s.first_name = 'John';
```

# SQL Queries:

## Joining together multiple tables

- Basic form  
**SELECT**  $\dots, r_i.A_j, \dots$   
**FROM**  $R_1 r_1, \dots, R_M r_M$   
**WHERE**  $\langle \text{condition} \rangle$
- When more than one relations in the **FROM** clause have an attribute named **A**, we refer to a specific **A** attribute as  $\langle \text{RelationName} \rangle.A$
- Hardest to get used to, yet most important feature of SQL

Produce a table that shows the pid, first name and last name of every student enrolled in the class with ID 1, along with the number of credit units in the "class 1" enrollment

```
SELECT s.pid, s.first_name,  
       s.last_name, e.credits  
FROM   students s, enrollment e  
WHERE  s.id = e.student  
       AND e.class = 1 ;
```

# (repeat)

**Classes**

id	name	number	date_code	start_time	end_time
1	Web stuff	MAS201	TuTh	2:00	3:20
2	Databases	CSE132A	TuTh	3:30	4:50
4	VLSI	CSE121	F	<i>null</i>	<i>null</i>

**Enrollment**

id	class	student	credits
1	1	1	4
2	1	2	3
3	4	3	4
4	1	3	3

**Students**

id	pid	first_name	last_name
1	8888888	John	Smith
2	1111111	Mary	Doe
3	2222222	<i>null</i>	Chen

# Take One: Understanding FROM as producing all combinations of tuples from the tables of the FROM clause

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE  s.id = e.student AND e.class = 1
```

“FROM” produces all 12 tuples made from one “students” tuple and one “enrollment” tuple

Student s part of the tuple				Enrollment e part of the tuple			
s.id	s.pid	s.first_name	s.last_name	e.id	e.class	e.student	e.credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	null	Chen	1	1	1	4
3	22..	null	Chen	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

## Take One: or understanding FROM as nested loops (producing all combinations)

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE s.id = e.student AND e.class = 1 ;
```

for **s** ranging over **students** tuples  
  for **e** ranging over **enrollment** tuples  
    output a tuple with all **s** attributes and **e** attributes

Student part of the tuple				Enrollment part of the tuple			
s.id	s.pid	s.first_name	s.last_name	e.id	e.class	e.student	e.credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	<i>null</i>	Chen	1	1	1	4
3	22..	<i>null</i>	Chen	2	1	2	3
3	22..	<i>null</i>	Chen	3	4	3	4
3	22..	<i>null</i>	Chen	4	1	3	3

# The order in FROM clause is unimportant

- FROM **students s, enrollment e**
- FROM **enrollment e, students s**

produce the same combinations (pairs) of student  
+ enrollment



## ... with shorter column names

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE  s.id = e.student AND e.class = 1 ;
```

“FROM” produces all 12 tuples made from one “students” tuple and one “enrollment” tuple

Student part of the tuple				Enrollment part of the tuple			
s.id	pid	first_name	last_name	e.id	class	student	credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	<i>null</i>	Chen	1	1	1	4
3	22..	<i>null</i>	Chen	2	1	2	3
3	22..	<i>null</i>	Chen	3	4	3	4
3	22..	<i>null</i>	Chen	4	1	3	3

# Understanding WHERE as qualifying the tuples that satisfy the condition

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE s.id = e.student AND e.class = 1 ;
```

s.id	s.pid	s.first_name	s.last_name	e.id	e.class	e.student	e.credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	null	Chen	1	1	1	4
3	22..	null	Chen	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

# Understanding SELECT as keeping the listed columns (highlighted below)

Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
<del>1</del>	<del>88..</del>	<del>John</del>	<del>Smith</del>	<del>2</del>	<del>1</del>	<del>2</del>	<del>3</del>
1	88..	John	Smith	3	4	3	4
<del>1</del>	<del>88..</del>	<del>John</del>	<del>Smith</del>	<del>4</del>	<del>1</del>	<del>3</del>	<del>3</del>
2	11..	Mary	Doe	1	1	1	4
<del>2</del>	<del>11..</del>	<del>Mary</del>	<del>Doe</del>	<del>2</del>	<del>1</del>	<del>2</del>	<del>3</del>
2	11..	Mary	Doe	3	4	3	4
<del>2</del>	<del>11..</del>	<del>Mary</del>	<del>Doe</del>	<del>4</del>	<del>1</del>	<del>3</del>	<del>3</del>
<del>3</del>	<del>22..</del>	<del>null</del>	<del>Chen</del>	<del>1</del>	<del>1</del>	<del>1</del>	<del>4</del>
<del>3</del>	<del>22..</del>	<del>null</del>	<del>Chen</del>	<del>2</del>	<del>1</del>	<del>2</del>	<del>3</del>
<del>3</del>	<del>22..</del>	<del>null</del>	<del>Chen</del>	<del>3</del>	<del>4</del>	<del>3</del>	<del>4</del>
<del>3</del>	<del>22..</del>	<del>null</del>	<del>Chen</del>	<del>4</del>	<del>1</del>	<del>3</del>	<del>3</del>

**SELECT s.pid, s.first\_name, s.last\_name, e.credits**

Students .pid	Students.first_name	Students.last_name	Enrollment.credits
88..	John	Smith	4
11..	Mary	Doe	3
22..	null	Chen	3

## Take Two on the previous exercises: The algebraic way

Produce a table that shows the pid, first name and last name of every student enrolled in the class with ID 1, along with the number of credit units in the “class 1” enrollment

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s JOIN enrollment e
       ON s.id = e.student
WHERE e.class = 1 ;
```

# Take two cont'd

## FROM clause result

Student part of the tuple				Enrollment part of the tuple			
s.id	pid	first_name	last_name	e.id	class	student	credits
1	88..	John	Smith	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
3	22..	<i>null</i>	Chen	3	4	3	4
3	22..	<i>null</i>	Chen	4	1	3	3

## WHERE clause result

s.id	pid	first_name	last_name	e.id	class	student	credits
1	88..	John	Smith	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
3	22..	<i>null</i>	Chen	3	4	3	4
<del>3</del>	<del>22..</del>	<del><i>null</i></del>	<del>Chen</del>	<del>4</del>	<del>1</del>	<del>3</del>	<del>3</del>

## Net result of the query is

s.pid	first_name	last_name	credits
88..	John	Smith	4
11..	Mary	Doe	3
22..	<i>null</i>	Chen	3

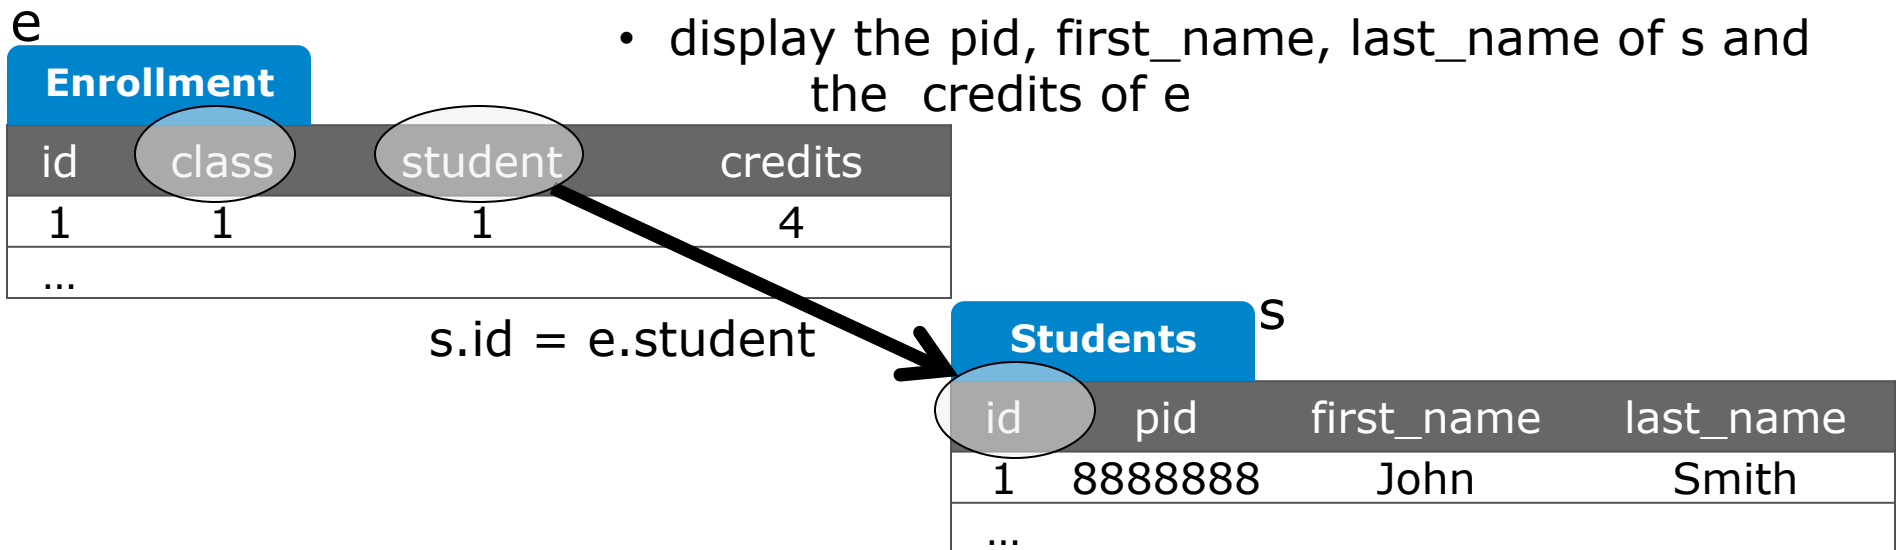
# Heuristics on writing queries

- Do you understand how queries work but have difficulty writing these queries yourself?
- The following heuristics will help you translate a requirement expressed in English into a query
  - The key point is to translate informal English into a precise English statement about which tuples your query should find in the database

## Hints for writing FROM/WHERE: Rephrase the statement, see it as a navigation across primary/foreign keys

Produce a table that shows the pid, first name and last name of every student enrolled in class 1, along with the number of credit units in his/her class 1 enrollment

- Find every enrollment tuple e
- that is an enrollment in class 1
  - i.e.,  $e.class = 1$
- and find the student tuple s that is connected to e
  - i.e., the student's id s.id appears in the enrollment tuple e as the foreign key e.student
- display the pid, first\_name, last\_name of s and the credits of e



- **Find every enrollment tuple e**

- that is an enrollment in class 1

- i.e., `e.class = 1`

- and find the student tuple s that is connected to e

- i.e., the student's id s.id appears in the enrollment tuple e as the foreign key e.student

- display the pid, first\_name, last\_name of s and the credits of e

FROM **enrollment e**

- Find every enrollment tuple e

- **that is an enrollment in class 1**

- **i.e., `e.class = 1`**

- and find the student tuple s that is connected to e

- i.e., the student's id s.id appears in the enrollment tuple e as the foreign key e.student

- display the pid, first\_name, last\_name of s and the credits of e

FROM enrollment e  
**WHERE e.class = 1**



- Find every enrollment tuple e
- that is an enrollment in class 1
  - i.e., e.class = 1
- **and find the student tuple s that is connected to e**
  - **i.e., the student's id s.id appears in the enrollment tuple e as the foreign key e.student**
- ~~display the pid, first\_name, last\_name of s and the~~

We could have also said "and find **every** student tuple s that is connected" but we used our knowledge that there is exactly one connected student and instead said "**the** student"

```
FROM enrollment e, students s
WHERE e.class = 1
      AND e.student = s.id
```

```
FROM enrollment e
      JOIN students s
      ON e.student = s.id
WHERE e.class = 1
```

- Find every enrollment tuple e
- that is an enrollment in class 1
  - i.e., e.class = 1
- and find the student tuple s that is connected to e
  - i.e., the student's id s.id appears in the enrollment tuple e as the foreign key e.student
- **display the pid, first\_name, last\_name of s and the credits of e**

```
SELECT s.pid, s.first_name, s.last_name,  
      e.credits
```

```
FROM enrollment e, students s
WHERE e.class = 1
      AND e.student = s.id
```

```
SELECT s.pid, s.first_name, s.last_name,  
      e.credits
```

```
FROM enrollment e
      JOIN students s
      ON e.student = s.id
WHERE e.class = 1
```

# SQL Queries: Nesting

- The **WHERE** clause can contain predicates of the form
  - `attr/value IN <query>`
  - `attr/value NOT IN <query>`
  - `attr/value = <query>`
- The predicate is satisfied if the `attr` or `value` appears in the result of the nested `<query>`
- Also
  - `EXISTS <query>`
  - `NOT EXISTS <query>`

# Nesting: Break the task into smaller

Produce a table that shows the pid, first name and last name of every student enrolled in the class named `MAS201`, along with the number of credit units in his/her `MAS201` enrollment

Note: We assume that there are no two classes with the same name

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE  e.class = (SELECT c.id
                  FROM classes c
                  WHERE c.number = 'MAS201'
                )
AND s.id = e.student
```

{[id:1]} -> 1

Nested queries modularize the task:  
Nested query finds the id of the MAS201 class.  
Then the outer query behaves as if there were a "1" in lieu of the subquery

# IN

Produce a table that shows the pid, first name and last name of every student enrolled in the class named `MAS201`, along with the number of credit units in his/her `MAS201` enrollment

Note: We assume that there are no two classes with the same name

```
SELECT s.pid, s.first_name, s.last_name, e.credits
FROM   students s, enrollment e
WHERE  e.class IN (SELECT c.id
                   FROM classes c
                   WHERE c.number = 'MAS201'
                  )
AND s.id = e.student
```



**{[id:1]}**

# Students + enrollments in class 1

## Vs Students who enrolled in class 1

Imagine a weird university where a student is allowed to enroll multiple times in the same class

Produce a table that shows the pid, first name and last name of every student enrolled in the class with ID 1, along with the number of credit units in the "class 1" enrollment

=> The same student may appear many times, once for each enrollment

```
SELECT s.pid, s.first_name,  
       s.last_name, e.credits  
FROM   students s, enrollment e  
WHERE  s.id = e.student  
       AND e.class = 1
```

Produce a table that shows the pid, first name and last name of every student who has enrolled at least once in the "class 1".

=> Each student will appear at most once

```
SELECT s.pid, s.first_name,  
       s.last_name  
FROM   students s  
WHERE  s.id IN ( SELECT e.student  
                  FROM enrollment e  
                  WHERE e.class = 1  
                )
```

# Uncorrelated subquery

```
SELECT s.pid, s.first_name,  
       s.last_name  
FROM   students s  
WHERE  s.id IN ( SELECT e.student  
                FROM enrollment e  
                WHERE e.class = 1  
              )
```

“Uncorrelated” in the sense  
that the nested query  
could be a standalone  
query

Some nested queries are  
correlated (example later)

# EXISTS

Display the students enrolled in class 1,  
only if the enrollment of class 2 is not empty

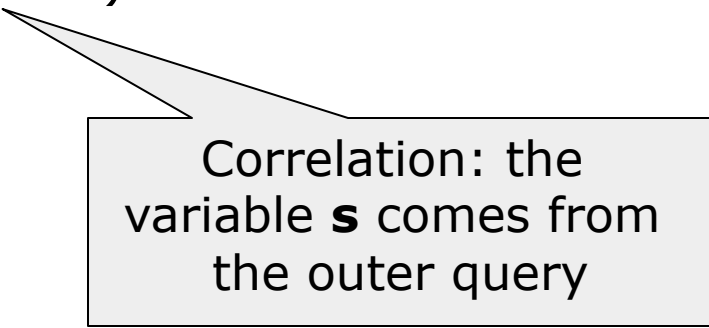
```
SELECT s.pid, s.first_name, s.last_name
FROM   students s
WHERE  s.id IN ( SELECT e.student
                  FROM enrollment e
                  WHERE e.class = 1
                )
AND EXISTS ( SELECT *
             FROM enrollment e
             WHERE e.class = 2
           )
```

Uncorrelated, also

# Correlated with EXISTS

Display the students enrolled in class 1

```
SELECT s.pid, s.first_name, s.last_name
FROM   students s
WHERE  EXISTS ( SELECT e.student
                  FROM enrollment e
                  WHERE e.class = 1
                        AND e.student = s.id )
```



Correlation: the variable **s** comes from the outer query



## Exercise, on thinking cardinalities of tuples in the results

```
SELECT  s.pid, s.first_name, s.last_name, e.credits
FROM    students s, enrollment e
WHERE   e.class IN (SELECT c.id
                    FROM classes c
                    WHERE c.number = 'MAS201'
                    )
        AND s.id = e.student
```

EXERCISE: Under what condition the above query always produces the same result with the query below?

```
SELECT  s.pid, s.first_name, s.last_name, e.credits
FROM    students s, enrollment e, classes c
WHERE   c.number = 'MAS201'
        AND s.id = e.student
        AND e.class = c.id
```

## Exercise: Multiple JOINS

Produce a table that shows the pid, first name and last name of every student enrolled in the MAS201 class along with the number of credit units in his/her 135 enrollment

Take One:

```
SELECT s.pid, s.first_name, s.last_name, e.credits  
FROM students s, enrollment e, classes c  
WHERE c.number = 'MAS201' AND s.id = e.student AND e.class = c.id
```

Take Two:

```
SELECT s.pid, s.first_name, s.last_name, e.credits  
FROM (students s JOIN enrollment e ON s.id = e.student)  
      JOIN classes c ON e.class = c.id  
WHERE c.number = 'MAS201'
```

## You can omit table names in SELECT, WHERE when attribute is unambiguous

```
SELECT  pid, first_name, last_name, credits
FROM    students, enrollment, classes
WHERE   number = 'MAS201'
        AND students.id = student
        AND class = classes.id ;
```

# SQL Queries, Aliases

- Use the same relation more than once in the same query or even the same **FROM** clause
- **Problem:** Find the Friday classes taken by students who take MAS201
  - also showing the students, i.e., produce a table where each row has the data of a MAS201 student and a Friday class he/she takes

Find the MAS201 students who take a Friday 11:00 am class

```
SELECT s.id, s.first_name, s.last_name, cF.number
FROM   students s, enrollment eF, classes cF
WHERE  date_code = 'F'
      AND eF.class = cF.id
      AND s.id = eF.student
      AND s.id IN
      (
        SELECT student
        FROM   enrollment e201, classes c201
        WHERE  c201.id = e201.class
              AND c201.number = 'MAS201'
      )
```

Nested query  
computes the id's of  
students enrolled in  
MAS201

# Multiple aliases may appear in the same FROM clause

Find the MAS201 students who take a Friday 11:00 am class

```
SELECT s.first_name, s.last_name, cF.number
FROM   students s, enrollment eF, classes cF,
        enrollment e201, classes c201
```

```
WHERE cF.date_code = 'F'
      AND eF.class = cF.id
      AND s.id = eF.student
      AND s.id = e201.student
      AND c201.id = e201.class
      AND c201.number = 'MAS201'
```

Under what conditions  
it computes the same  
result with previous  
page?

# DISTINCT

Find the other classes taken by MAS201 students  
(I don't care which students)

```
SELECT DISTINCT cOther.number
FROM   enrollment eOther, classes cOther,
       enrollment e201, classes c201
WHERE  eOther.class = cOther.id
       AND eOther.student = e201.student
       AND c201.id = e201.class
       AND c201.number = 'MAS201'
```

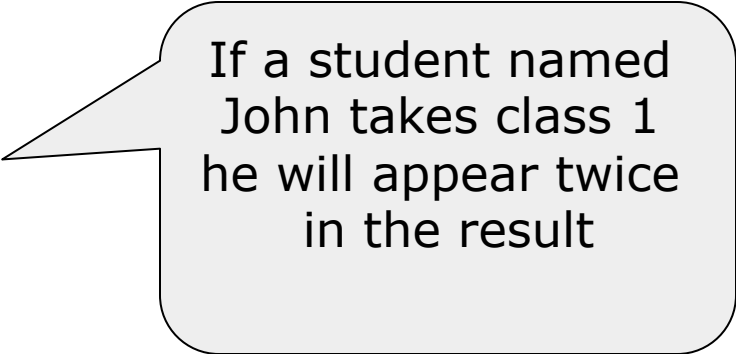
# UNION ALL

Find all student ids for students who have taken class 1 or are named 'John'

```
( SELECT e.student  
  FROM enrollment e  
  WHERE e.class=1 )
```

## **UNION ALL**

```
( SELECT s.id AS student  
  FROM student s  
  WHERE s.first_name='John'  
)
```



If a student named John takes class 1 he will appear twice in the result



# UNION with non –duplicate results

```
( SELECT e.student  
  FROM enrollment e  
  WHERE e.class=1 )  
UNION  
( SELECT s.id AS student  
  FROM student s  
  WHERE s.first_name='John'  
)
```

# SQL Queries: Aggregation & Grouping

- Aggregate functions: SUM, AVG, COUNT, MIN, MAX, and recently user defined functions as well
- GROUP BY

**Employee**

Name	Dept	Salary
Joe	Toys	45
Nick	PCs	50
Jim	Toys	35
Jack	PCs	40

**Example:** Find the average salary of all employees:

```
SELECT AVG(Salary) AS AvgSal  
FROM Employee
```

AvgSal
42.5

**Example:** Find the average salary for each department:

```
SELECT Dept, AVG(Salary) AS AvgSal  
FROM Employee  
GROUP BY Dept
```

Dept	AvgSal
Toys	40
PCs	45

# SQL Grouping:

## Conditions that Apply on Groups

- **HAVING** <condition> may follow a **GROUP BY** clause
- If so, the condition applies to each group, and groups not satisfying the condition are eliminated
- **Example:** Find the average salary in each department that has more than 1 employee:

```
SELECT Dept, AVG(Salary) AS AvgSal
FROM Employee
GROUP BY Dept
HAVING COUNT(Name) >1
```

## Let's mix features we've seen: Aggregation after joining tables

- **Problem:** List all enrolled students and the number of total credits for which they have registered

```
SELECT  students.id, first_name, last_name, SUM(credits)
FROM    students, enrollment
WHERE   students.id = enrollment.student
GROUP BY students.id, first_name, last_name
```

# ORDER BY and LIMIT

Order the student id's of class 2 students according to their class 2 credits, descending

```
SELECT e.student  
FROM enrollment e  
WHERE e.class = 2  
ORDER BY e.credits DESC
```

Order the student id's of class 2 students according to their class 2 credits, descending **and display the Top 10**

```
SELECT e.student  
FROM enrollment e  
WHERE e.class = 2  
ORDER BY e.credits DESC  
LIMIT 10
```

# Combining features

Find the Top-5 classes taken by students of class 2, i.e., the 5 classes (excluding class 2 itself) with the highest enrollment of class 2 students, display their numbers and how many class 2 students they have

```
SELECT cF.number, COUNT(*)
FROM   enrollment eF, classes cF
WHERE  eF.class = cF.id
      AND NOT(eF.class = 2)
      AND eF.student IN
      (
        SELECT student
        FROM   enrollment e2
        WHERE  e201.class = 2
      )
GROUP BY cF.id, cF.number
ORDER BY cF.number
LIMIT 5
```

Grouping by both id and number ensures correctness even if multiple classes have the same number

# The outerjoin operator

- New construct in FROM clause
- R LEFT OUTER JOIN S ON R.<attr of R>=S.<attr of J>
- R FULL OUTER JOIN S ON R.<attr of R>=S.<attr of J>

R		S	
RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	3	SV3

SELECT \*  
FROM R LEFT OUTERJOIN S ON R.RJ=S.SJ

RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	Null	Null

SELECT \*  
FROM R FULL OUTERJOIN S ON R.RJ=S.SJ

RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	Null	Null
Null	Null	3	SV3

# An application of outerjoin

- **Problem:** List all students and the number of total credits for which they have registered
  - Notice that you must also list non-enrolled students

```
SELECT  students.id, first_name, last_name, SUM(credits)
FROM    students LEFT OUTER JOIN enrollment ON
        students.id = enrollment.student
GROUP BY students.id, first_name, last_name
```



# SQL: More Bells and Whistles ...

- Pattern matching conditions
  - `<attr> LIKE <pattern>`

Retrieve all students whose name contains "Sm"

```
SELECT *  
FROM Students  
WHERE name LIKE '%Sm%'
```

# ...and a Few “Dirty” Points

- **Null values**

- All comparisons involving NULL are **false** by definition
- All aggregation operations, except **COUNT (\*)**, ignore NULL values

# Null Values and Aggregates

- Example:

R	
a	b
x	1
x	2
x	null
null	null
null	null

```
SELECT COUNT (a) , COUNT (b) , AVG (b) , COUNT (*)  
FROM R  
GROUP BY a
```

count(a)	count(b)	avg(b)	count(*)
3	2	1.5	3
0	0	null	2

# Universal Quantification by Negation (difficult)

Problem:

- Find the students that take **every** class 'John Smith' takes

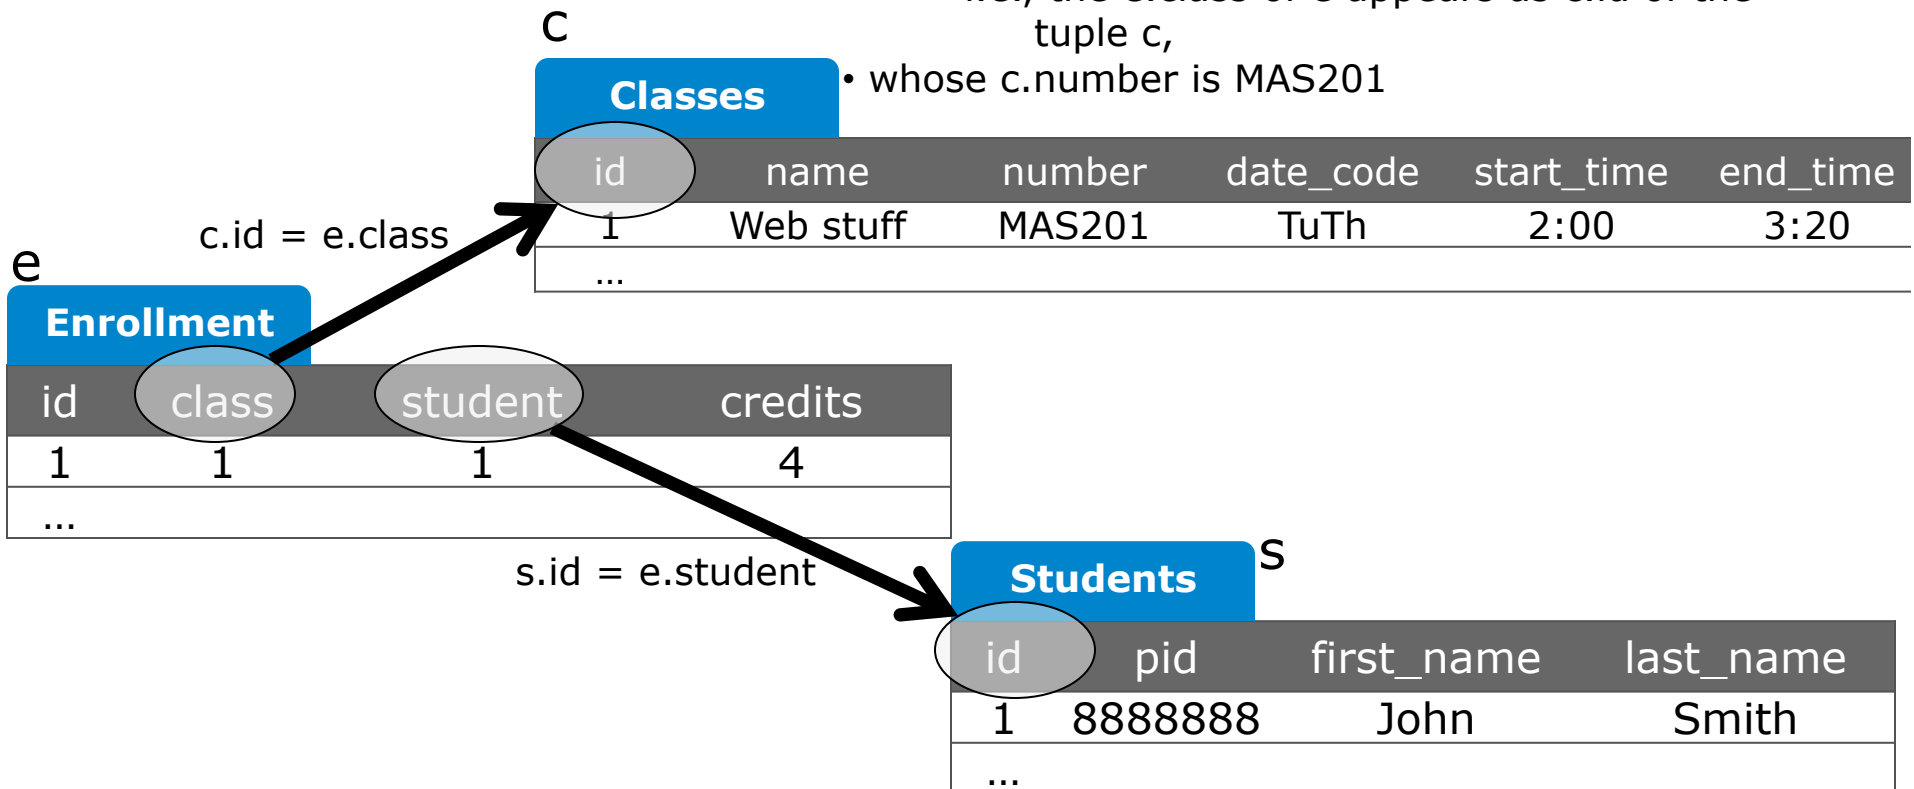
Rephrase:

- Find the students such that there is no class that 'John Smith' takes and they do not take

## Hints for writing FROM/WHERE: Rephrase the statement, see it as a navigation across primary/foreign keys

Produce a table that shows the pid, first name and last name of every student enrolled in the MAS201 class along with the number of credit units in his/her 135 enrollment

- Find every combination of a students tuple s, an enrollment tuple e, c where
- the students tuple s,
- is connected to the enrollment tuple e
  - i.e., the s.id appears in the enrollment tuple e as e.student,
- and e is connected to the classes tuple c
  - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is MAS201



- Find any students tuple s,
- that is connected to an enrollment tuple e
  - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
  - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is MAS201

FROM students AS s

- Find any students tuple s,
- that is connected to an enrollment tuple e
  - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
  - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is MAS201

*Take One: Declarative*

FROM students AS s,  
enrollment AS e  
WHERE s.id = e.student

*Take Two: Algebraic*

FROM students AS s  
JOIN enrollment AS e  
ON s.id = e.student

- 
- Find any students tuple s,
  - that is connected to an enrollment tuple e
    - i.e., whose s.id appears in an enrollment tuple e as e.student,
  - and e is connected to a classes tuple c
    - i.e., the e.class of e appears as c.id of the tuple c,
  - whose c.number is MAS201

*Take One: Declarative*

```
FROM students AS s,  
      enrollment AS e,  
      classes AS c  
WHERE s.id = e.student  
      AND c.id = e.class
```

*Take Two: Algebraic*

```
FROM ( students AS s  
      JOIN enrollment AS e  
      ON s.id = e.student )  
JOIN classes AS c  
ON c.id = e.class
```

- 
- Find any students tuple s,
  - that is connected to an enrollment tuple e
    - i.e., whose s.id appears in an enrollment tuple e as e.student,
  - and e is connected to a classes tuple c
    - i.e., the e.class of e appears as c.id of the tuple c,
  - whose c.number is MAS201

*Take One: Declarative*

```
FROM students AS s,  
      enrollment AS e,  
      classes AS c  
WHERE s.id = e.student  
      AND c.id = e.class  
      AND c.number = 'MAS201'
```

*Take Two: Algebraic*

```
FROM ( students AS s  
      JOIN enrollment AS e  
      ON s.id = e.student )  
      JOIN classes AS c  
      ON c.id = e.class  
WHERE c.number = 'MAS201'
```



- Find any students tuple s,
- that is connected to an enrollment tuple e
  - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
  - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is MAS201

FROM students AS s

- Find any students tuple s,
- that is connected to an enrollment tuple e
  - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
  - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is MAS201

*Take One: Declarative*

FROM students AS s,  
 enrollment AS e  
 WHERE s.id = e.student

*Take Two: Algebraic*

FROM students AS s,  
 JOIN enrollment AS e  
 ON s.id = e.student

# Breaking a query into pieces with WITH

Find the 5 classes whose students have the busiest courseload, i.e., the 5 classes whose students have the highest average of quarter credits

Defines a table  
"courseload" that lives for  
the duration of this query  
only

```
WITH courseload AS  
( SELECT e.student, SUM(credits) AS total_credits  
  FROM enrollment e  
   GROUP BY e.student )  
SELECT e.class, AVG(c.total_credits)  
FROM enrollment e, courseload c  
WHERE e.student = c.student  
GROUP BY e.class  
ORDER BY AVG(c.total_credits) DESC  
LIMIT 5
```

# Avoid repeating aggregates

```
WITH courseload AS
( SELECT e.student, SUM(credits) AS total_credits
  FROM enrollment e
  GROUP BY e.student )
SELECT e.class, AVG(c.total_credits)
FROM enrollment e, courseload c
WHERE e.student = c.student
GROUP BY e.class
ORDER BY AVG(c.total_credits) DESC
LIMIT 5
```

← Equivalent  
↓

```
WITH courseload AS
( SELECT e.student, SUM(credits) AS total_credits
  FROM enrollment e
  GROUP BY e.student )
SELECT e.class, AVG(c.total_credits) AS credits_avg
FROM enrollment e, courseload c
WHERE e.student = c.student
GROUP BY e.class
ORDER BY credits_avg DESC
LIMIT 5
```

# Breaking a query into pieces with nesting in the FROM clause

Find the 5 classes whose students have the busiest courseload, i.e., the 5 classes whose students have the highest average of quarter credits

Also defines a table, identical to the "courseload" except it has no name

```
SELECT e.class, AVG(c.total_credits) AS credits_avg
FROM enrollment e,
  ( SELECT e.student, SUM(credits) AS total_credits
    FROM enrollment e
    GROUP BY e.student ) c
WHERE e.student = c.student
GROUP BY e.class
ORDER BY credits_avg DESC
LIMIT 5
```

## and nesting in the **SELECT** clause

Find the 5 classes whose students have the busiest courseload, i.e., the 5 classes whose students have the highest average of quarter credits

```
SELECT e.class, AVG(  
    ( SELECT SUM(es.credits)  
      FROM enrollment es  
      WHERE es.student = e.student )  
    ) AS credits_avg  
FROM enrollment e  
GROUP BY e.class  
ORDER BY credits_avg DESC  
LIMIT 5
```

## Discussed in class and discussion section

How to solve in easy steps the following complex query:

Create a table that shows all time slots (date, start time, end time) when students of MAS201 attend a lecture of another class;  
Show also how many students attend a class at each time slot.

# SQL as a Data Manipulation Language: Insertions

- Inserting tuples

```
INSERT INTO R (A1, ..., Ak)  
VALUES (v1, ..., vk) ;
```

- Some values may be left NULL
- Use results of queries for insertion

```
INSERT INTO R  
SELECT ...  
FROM ...  
WHERE ...
```

- Insert in Students 'John Doe' with A# 99999999

```
INSERT INTO students  
(pid, first_name, last_name)  
VALUES  
(`99999999`, `John`, `Doe`)
```

- Enroll all MAS201 students into CSE132A

```
INSERT INTO enrollment (class,  
student)  
SELECT c132a.id, student  
FROM classes AS c135, enrollment,  
classes AS c132a  
WHERE c135.number='MAS201' AND  
enrollment.class=c135.id AND  
c132a.number='CSE132A'
```

# SQL as a Data Manipulation Language: Updates and Deletions

- Deletion basic form: delete every tuple that satisfies **<cond>**:

**DELETE FROM R**

**WHERE <cond>**

- Update basic form: update every tuple that satisfies **<cond>** in the way specified by the **SET** clause:

**UPDATE R**

**SET  $A_1 = \langle \text{exp}_1 \rangle, \dots, A_k = \langle \text{exp}_k \rangle$**

**WHERE <cond>**

- Delete "John" "Smith"
- DELETE FROM students  
WHERE  
first\_name='John' AND  
last\_name='Smith'

- Update the registered credits of all MAS201 students to 5

UPDATE enrollment  
SET credits=5  
WHERE class=1

UPDATE enrollment  
SET credits=5  
WHERE class IN  
(SELECT id FROM classes  
WHERE number='MAS201')