

DSE210 Worksheet 6

Orysya Stas

1. $\pi(h) = \frac{3}{4}$

given h

$\Pr(a \text{ lot} | h) = \frac{2}{3}$

$\Pr(\text{little} | h) = \frac{1}{6}$

$\Pr(\text{silent} | h) = \frac{1}{6}$

$\pi(s) = \frac{1}{4}$

given s

$\Pr(a \text{ lot} | s) = \frac{1}{6}$

$\Pr(\text{little} | s) = \frac{1}{6}$

$\Pr(\text{silent} | s) = \frac{2}{3}$

a. $P(s|\text{little}) = \frac{P(\text{little}|s)P(s)}{P(\text{little})} = \frac{P(\text{little}|s)P(s)}{P(\text{little}|s)P(s) + P(\text{little}|h)P(h)} = \frac{(1/6)(1/4)}{(1/6)(1/4) + (1/6)(3/4)} = \frac{1}{4}$

$P(h|\text{little}) = \frac{P(\text{little}|h)P(h)}{P(\text{little})} = \frac{(1/6)(3/4)}{4/24} = \frac{3}{4}$

↳ he is most likely happy

b. IP of prediction in part (a) being incorrect

$P(s|\text{little}) = \frac{1}{4}$

2. $X = [-1, 1]$ $Y = \{1, 2, 3\}$ $\pi_1 = \frac{1}{3}$ $\pi_2 = \frac{1}{6}$ $\pi_3 = \frac{1}{2}$ P_1, P_2, P_3 given

$h^*(x) = \arg \max_j \pi_j \cdot P_j(x) = \begin{cases} 1 & \text{if } -1 \leq x \leq 0 \\ 3 & \text{if } 0 < x < 1 \end{cases}$

3a. positively correlated

b. positively correlated

c. negatively correlated

4. $\text{corr}(\text{wife age}, \text{husband age}) = 1$

5. give the parameters of the (unique) bivariate Gaussian that satisfies these props

a. $\mu_x = 2$ $\text{std}(x) = 1$ $\mu_y = 2$ $\text{std}(y) = 0.5$ $\text{corr}(x, y) = -0.5$

$\text{corr}(x, y) = \frac{\text{cov}(X, Y)}{\text{std}(x)\text{std}(y)} \quad -0.5 = \frac{\text{cov}(X, Y)}{1 \cdot 0.5} \quad \text{cov}(X, Y) = -0.25$

covariance matrix $\Sigma = \begin{pmatrix} \text{var}(X) & \text{cov}(X, Y) \\ \text{cov}(X, Y) & \text{var}(Y) \end{pmatrix} = \begin{pmatrix} 1 & -0.25 \\ -0.25 & 0.25 \end{pmatrix} \quad \mu = \begin{pmatrix} 2 \\ 2 \end{pmatrix}$

b. $\mu_x = 1 = \mu_y$ $\text{std}(x) = \text{std}(y) = 1$


$\mu = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$ $\text{cov}(X, Y) = 1$

covariance matrix $\Sigma = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$

6. sketch shapes of Gaussians $N(\mu, \Sigma)$


a. $\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$

$\Sigma = \begin{pmatrix} 9 & 0 \\ 0 & 1 \end{pmatrix}$



b. $\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$

$\Sigma = \begin{pmatrix} 1 & -0.75 \\ -0.75 & 1 \end{pmatrix}$



Worksheet 6

February 18, 2017

1 Worksheet 6 – Generative models 2

1.1 Orysyia Stus

1.2 2.19.2017

<https://github.com/mas-dse/ostus/tree/master/DSE210>

```
In [1]: import pandas as pd
import numpy as np
import math
import sklearn
from sklearn.naive_bayes import MultinomialNB
from sklearn import metrics
from sklearn import datasets
from scipy.stats import multivariate_normal
import random
import seaborn as sns
import matplotlib.pyplot as plt
import re

%matplotlib inline
```

C:\Users\Orysyia\Anaconda\lib\site-packages\IPython\html.py:14: ShimWarning: The `IPython.html.widgets` has moved to `ipywidgets`.", ShimWarning)

1.3 Problem 7

For each of the two Gaussians in the previous problem, check your answer using Python: draw 100 random samples from that Gaussian and plot it.


Using: https://docs.scipy.org/doc/numpy/reference/generated/numpy.random.multivariate_normal.html

1.3.1 7a

```
In [2]: mean = [0,0]
cov = [[9,0], [0,1]]
x, y = np.random.multivariate_normal(mean, cov, 100).T
plt.scatter(x, y)
```

```
plt.axis('equal')
plt.title('Problem 7a Gaussian')
```


Out[2]: <matplotlib.text.Text at 0xbee0550>



1.3.2 7b

```
In [3]: mean = [0,0]
cov = [[1, -0.75], [-0.75, 1]]
x, y = np.random.multivariate_normal(mean, cov, 100).T
plt.scatter(x, y)
plt.axis('equal')
plt.title('Problem 7b Gaussian')
```

Out[3]: <matplotlib.text.Text at 0xc0627b8>




1.4 Problem 8

Consider the linear classifier. Sketch the decision boundary in \mathbb{R}^2 . Make sure to label precisely where the boundary intersects the coordinate axes, and also indicate which side of the boundary is the positive side.

```
In [4]: x = np.linspace(-10, 4, 10, endpoint=True)
        y = (3 * x + 12) / 4
        fig = plt.figure()
        ax = fig.add_subplot(111)
        ax.plot(x, y)
        ax.spines['right'].set_color('none')
        ax.spines['top'].set_color('none')
        ax.xaxis.set_ticks_position('bottom')
        ax.spines['bottom'].set_position(('data', 0))
        ax.yaxis.set_ticks_position('left')
        plt.annotate(r'positive', xy=(-6, 2), xycoords='data', xytext=(+50, +30), t
```

```
Out[4]: <matplotlib.text.Annotation at 0xc355080>
```



2 Problem 9: Handwritten digit recognition using a Gaussian generative model

Handwritten digit recognition using a Gaussian generative model. In class, we mentioned the MNIST data set of handwritten digits. You can obtain it from: <http://yann.lecun.com/exdb/mnist/index.html>

In this problem, you will build a classifier for this data, by modeling each class as a multivariate (784-dimensional) Gaussian.

Refer to the following link for review:

<http://www.eggie5.com/68-mnist-gaussian-classifier>

2.1 Part (a)

Upon downloading the data, you should have two training files (one with images, one with labels) and two test files. Unzip them.

In order to load the data into Python you will find the following code helpful:

<http://cseweb.ucsd.edu/~dasgupta/dse210/loader.py>

For instance, to load in the training data, you can use:

```
x,y = loadmnist('train-images-idx3-ubyte', 'train-labels-idx1-ubyte')
```

This will set x to a 60000×784 array where each row corresponds to an image, and y to a length-60000 array where each entry is a label (0-9). There is also a routine to display images: use `displaychar(x[0])` to show the first data point, for instance.

```
In [5]: import loader as loader
        x, y = loader.loadmnist('train-images.idx3-ubyte', 'train-labels.idx1-ubyte')
        print 'The shape of x is:', x.shape
        print 'The shape of y is:', y.shape
```

The shape of x is: (60000L, 784L)

The shape of y is: (60000L,)

2.1.1 Examine the digit data (as an array and image)

```
In [6]: print x[1]
        loader.displaychar(x[1])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 51 159 253 159 50  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 48 238 252 252 252 237  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 54 227 253 252 239 233 252 57  6  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0 10 60 224 252 253 252 202 84 252
253 122  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 163 252 252 252 253 252 252 96 189 253 167  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 51 238 253 253 190 114 253 228
47 79 255 168  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 48 238 252 252 179 12 75 121 21  0  0 253 243 50  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 38 165 253 233 208 84  0  0
  0  0  0  0 253 252 165  0  0  0  0  0  0  0  0  0  0  0
  0  7 178 252 240 71 19 28  0  0  0  0  0  0 253 252 195  0
  0  0  0  0  0  0  0  0  0  0  0  0 57 252 252 63  0  0  0
  0  0  0  0  0  0 253 252 195  0  0  0  0  0  0  0  0  0
  0  0  0 198 253 190  0  0  0  0  0  0  0  0  0  0  0 255 253
196  0  0  0  0  0  0  0  0  0  0  0  0 76 246 252 112  0  0
  0  0  0  0  0  0  0  0 253 252 148  0  0  0  0  0  0  0
  0  0  0  0 85 252 230 25  0  0  0  0  0  0  0  0  7 135
253 186 12  0  0  0  0  0  0  0  0  0  0  0  0 85 252 223  0
  0  0  0  0  0  0  0  7 131 252 225 71  0  0  0  0  0  0
  0  0  0  0  0  0 85 252 145  0  0  0  0  0  0  0 48 165
252 173  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 86 253
225  0  0  0  0  0  0  0 114 238 253 162  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 85 252 249 146 48 29 85 178 225 253
223 167 56  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 85 252 252 252 229 215 252 252 252 196 130  0  0  0  0  0  0  0
```

```

0 0 0 0 0 0 0 0 0 0 28 199 252 252 253 252 252 233
145 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 25 128 252 253 252 141 37 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]

```



2.2 Part (b)

Split the training set into two pieces - a training set of size 50000, and a separate validation set of size 10000. Also load in the test data.

```

In [7]: from sklearn.cross_validation import train_test_split
        X_train, X_validation, y_train, y_validation = train_test_split(x, y, test_
        print 'The shape of the x train data set is: ', X_train.shape
        print 'The shape of the y train data set is: ', y_train.shape
        print 'The shape of the x validation data set is: ', X_validation.shape
        print 'The shape of the y validation data set is: ', y_validation.shape

```

The shape of the x train data set is: (50000L, 784L)

The shape of the y train data set is: (50000L,)

```
The shape of the x validation data set is: (10000L, 784L)
The shape of the y validation data set is: (10000L,)
```

```
C:\Users\Orysy\Anaconda\lib\site-packages\sklearn\cross_validation.py:44: Deprecat
    "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [8]: X_test, y_test = loader.loadmnist('t10k-images.idx3-ubyte', 't10k-labels.idx1-ubyte')
        print 'The shape of the x test data set is: ', X_test.shape
        print 'The shape of the y test data set is: ', y_test.shape
```

```
The shape of the x test data set is: (10000L, 784L)
The shape of the y test data set is: (10000L,)
```

2.3 Part (c)

Now fit a Gaussian generative model to the training data of 50000 points:

Determine the class probabilities: what fraction of π_0 of the training points are

Fit a Gaussian to each digit, by finding the mean and the covariance of the corresponding data. Let the Gaussian for the j th digit by $P_j = N(\mu_j, \sigma_j)$.

Using these two pieces of information, you can classify new images x using Bayes' rule: simply pick the digit j for which $\pi_j P_j(x)$ is largest.

```
In [9]: from sklearn.naive_bayes import GaussianNB
        clf = GaussianNB()
        clf.fit(X_train, y_train)
```

```
Out[9]: GaussianNB(priors=None)
```

2.3.1 Show the fraction (percentage) of images that belong to each class. Associated with the π_j or the prior probabilities. The class distribution is fairly uniform.

```
In [10]: priors = clf.class_prior_
         length = len(priors)
         for i in range(length):
             percentage = (priors[i])
             print 'The prior probability of class image %s is' %i, '{0:.4f}'.format(percentage)
```

```
The prior probability of class image 0 is 0.0994
The prior probability of class image 1 is 0.1118
The prior probability of class image 2 is 0.0996
The prior probability of class image 3 is 0.1020
The prior probability of class image 4 is 0.0972
The prior probability of class image 5 is 0.0901
```


The prior probability of class image 6 is 0.0981
The prior probability of class image 7 is 0.1042
The prior probability of class image 8 is 0.0978
The prior probability of class image 9 is 0.0997

2.3.2 Fit the train data set to the model used above (note: the validation set was not yet used/considered)

```
In [11]: classes = clf.classes_
         posteriors = []

         def class_grouping(class_id):
             grouping = []
             for i, group in enumerate(X_train):
                 if y_train[i] == class_id:
                     grouping.append(group)
             grouping = np.matrix(grouping)
             return grouping

         posteriors = []
         for c in classes:
             grouping = class_grouping(c)
             mean = np.array(grouping.mean(0))[0]
             cov = np.cov(grouping.T)
             Px = multivariate_normal(mean, cov, allow_singular=True)
             posteriors.append(Px)
```

2.3.3 Examining model predictions based on Bayes probability

```
In [12]: x = random.choice(X_test)
         actual = y_test[x]

         bayes_prob = []
         for c in classes:
             prob = [c, priors[c] * posteriors[c].pdf(x)]
             formatting_function = np.vectorize(lambda f: format(f, '6.3E'))
             bayes_prob.append(prob)

         print 'The Bayes probability found is \n', bayes_prob
         prediction = max(bayes_prob, key= lambda a: a[1])
```

The Bayes probability found is

[[0, 0.0], [1, 0.0], [2, 0.0], [3, 0.0], [4, 0.0], [5, 0.0], [6, 0.0], [7, 0.0], [8, 0.0], [9, 0.0]]

2.3.4 Building the classifier

```
In [13]: Y = []
        for x in X_test:
            bayes_prob = []
            for c in classes:
                prob = [c, priors[c] * posteriors[c].pdf(x)]
                bayes_prob.append(prob)
            prediction = max(bayes_prob, key= lambda a: a[1])
            Y.append(prediction[0])

In [14]: errors = (y_test != Y).sum()
        total = X_test.shape[0]
        print("Error rate: %d/%d = %f" % (errors, total, (errors/float(total))))

Error rate: 9020/10000 = 0.902000
```


2.3.5 Our naively implemented Gaussian Classifier achieved a 9.8% success rate. Note such a low success rate can be due to using 'pdf' vs 'logpdf' which results in very small probabilities. Loss in probability precision results in inaccurate predictions.

2.3.6 Comparing 100 of the predicted and actual classes for the digit data, the model predicted all of the digits to be 0s.

```
In [15]: def displayimage(image):
        plt.imshow(np.reshape(image, (28, 28)), cmap=plt.cm.gray)
        plt.axis('off')

        indices = np.array(np.where((y_test != Y)==True))[0]
        indices = indices[0:100]
        index = 0
        rows = len(indices)/10
        cols = 10

        plt.figure(figsize=(30,15))
        for i in indices:
            index += 1
            plt.subplot(rows, cols, index)
            plt.subplots_adjust(hspace=.5)
            displayimage(X_test[i])
            plt.title('pre:%i, act:%i' % (Y[i], y_test[i]), fontsize = 20)
```



2.3.7 Building the classifier: Using log probabilities

```
In [16]: Y = []
         for x in X_test:
             bayes_prob = []
             for c in classes:
                 prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
                 bayes_prob.append(prob)
             prediction = max(bayes_prob, key=lambda a: a[1])
             Y.append(prediction[0])

In [17]: errors = (y_test != Y).sum()
         total = X_test.shape[0]
         print("Error rate: %d/%d = %f" % (errors, total, (errors/float(total))))
```

Error rate: 1862/10000 = 0.186200

2.3.8 Our naively implemented Gaussian Classifier achieved a 81.4% success rate. Using log-pdf resulted in more precise probabilities, thus a more accurate model.

2.3.9 Comparing 100 of the predicted and actual classes for the digit data, the model predicted all of the digits to be 0s.

```
In [18]: indices = np.array(np.where((y_test != Y)==True))[0]
         indices = indices[0:100]
         index = 0
         rows = len(indices)/10
```


```
plt.figure(figsize=(30,15))
for i in indices:
    index += 1
    plt.subplot(rows, cols, index)
    plt.subplots_adjust(hspace=.5)
    displayimage(X_test[i])
    plt.title('pre:%i, act:%i' % ( Y[i], y_test[i]), fontsize = 20)
```



The model (no validation set used) has an accuracy of 0.8138

11

5	0.92	0.55	0.69	892
6	0.94	0.92	0.93	958
7	0.93	0.63	0.75	1028
8	0.56	0.88	0.68	974
9	0.64	0.90	0.75	1009
avg / total	0.85	0.81	0.82	10000



2.3.10 The model above misclassified digits 5 & 7 most often.

2.4 Part (d)

One last step is needed: it is important to smooth the covariance matrices, and the usual way to do this is to add in cI , where c is some constant and I is the identity matrix. What value of c is right? Use the validation set to help you choose. That is, choose the value of c for which the resulting classifier makes the fewest mistakes on the validation set. What value of c did you get?

```
In [20]: smoothing_c = range(0, 5000, 500)
         error_rates = []
         for sc in smoothing_c:
             posteriors = []
             for c in classes:
                 grouping = class_grouping(c)
                 mean = np.array(grouping.mean(0))[0]
                 cov = np.cov(grouping, rowvar=0)
                 cov_smoothed = cov + (sc * np.eye(mean.shape[0]))
                 p_x = multivariate_normal(mean, cov_smoothed, allow_singular=True)
                 posteriors.append(p_x)

             Y = []
             for x in X_validation:
                 bayes_prob = []
                 for c in classes:
                     prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
                     bayes_prob.append(prob)
                 prediction = max(bayes_prob, key= lambda a: a[1])
                 Y.append(prediction[0])

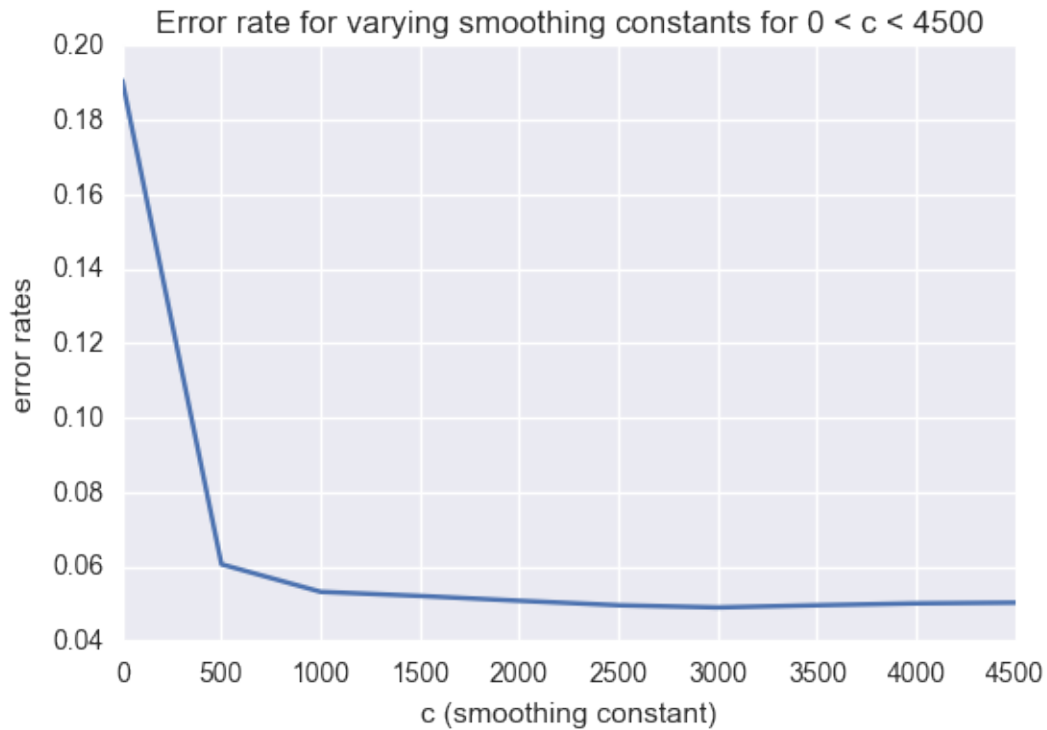
             errors = (y_validation != Y).sum()
             total = X_validation.shape[0]
             error_rate = errors/float(total)
             error_rates.append(error_rate)
             print("Error rate for c= %s: %d/%d = %f" % (sc, errors, total, error_rate))

Error rate for c= 0: 1906/10000 = 0.190600
Error rate for c= 500: 606/10000 = 0.060600
Error rate for c= 1000: 532/10000 = 0.053200
Error rate for c= 1500: 521/10000 = 0.052100
Error rate for c= 2000: 508/10000 = 0.050800
Error rate for c= 2500: 496/10000 = 0.049600
Error rate for c= 3000: 490/10000 = 0.049000
Error rate for c= 3500: 496/10000 = 0.049600
Error rate for c= 4000: 501/10000 = 0.050100
Error rate for c= 4500: 503/10000 = 0.050300

In [21]: plt.plot(smoothing_c, error_rates)
         plt.xlabel('c (smoothing constant)')
```

```
plt.ylabel('error rates')
plt.title('Error rate for varying smoothing constants for  $0 < c < 4500$ ')
```

Out[21]: <matplotlib.text.Text at 0x220280f0>



2.4.1 The optimal c lies within $2500 < \text{smoothing_c} < 3500$, iterative smoothing is run again to determine the optimal smoothing c .

```
In [22]: smoothing_c = range(2500, 3500, 100)
error_rates = []
for sc in smoothing_c:
    posteriors = []
    for c in classes:
        grouping = class_grouping(c)
        mean = np.array(grouping.mean(0))[0]
        cov = np.cov(grouping, rowvar=0)
        cov_smoothed = cov + (sc * np.eye(mean.shape[0]))
        p_x = multivariate_normal(mean, cov_smoothed, allow_singular=True)
        posteriors.append(p_x)

Y = []
for x in X_validation:
    bayes_prob = []
```

```

    for c in classes:
        prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
        bayes_prob.append(prob)
    prediction = max(bayes_prob, key= lambda a: a[1])
    Y.append(prediction[0])

errors = (y_validation != Y).sum()
total = X_validation.shape[0]
error_rate = errors/float(total)
error_rates.append(error_rate)
print("Error rate for c= %s: %d/%d = %f" % (sc, errors, total, error_rate))

```

```

Error rate for c= 2500: 496/10000 = 0.049600
Error rate for c= 2600: 490/10000 = 0.049000
Error rate for c= 2700: 489/10000 = 0.048900
Error rate for c= 2800: 490/10000 = 0.049000
Error rate for c= 2900: 488/10000 = 0.048800
Error rate for c= 3000: 490/10000 = 0.049000
Error rate for c= 3100: 491/10000 = 0.049100
Error rate for c= 3200: 493/10000 = 0.049300
Error rate for c= 3300: 495/10000 = 0.049500
Error rate for c= 3400: 495/10000 = 0.049500

```

```

In [23]: plt.plot(smoothing_c, error_rates)
         plt.xlabel('c (smoothing constant)')
         plt.ylabel('error rates')
         plt.title('Error rate for varying smoothing constants for 2500 < c < 3400')

```

```

Out[23]: <matplotlib.text.Text at 0x200b4d30>

```




2.4.2 The optimal c lies within $2800 < \text{smoothing_c} < 3000$, iterative smoothing is run again to determine the optimal smoothing c .

```
In [24]: smoothing_c = range(2800, 3000, 10)
error_rates = []
for sc in smoothing_c:
    posteriors = []
    for c in classes:
        grouping = class_grouping(c)
        mean = np.array(grouping.mean(0))[0]
        cov = np.cov(grouping, rowvar=0)
        cov_smoothed = cov + (sc * np.eye(mean.shape[0]))
        p_x = multivariate_normal(mean, cov_smoothed, allow_singular=True)
        posteriors.append(p_x)

Y = []
for x in X_validation:
    bayes_prob = []
    for c in classes:
        prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
        bayes_prob.append(prob)
    prediction = max(bayes_prob, key= lambda a: a[1])
    Y.append(prediction[0])
```

```


errors = (y_validation != Y).sum()
total = X_validation.shape[0]
error_rate = errors/float(total)
error_rates.append(error_rate)
print("Error rate for c= %s: %d/%d = %f" % (sc, errors, total, error_r

Error rate for c= 2800: 490/10000 = 0.049000
Error rate for c= 2810: 490/10000 = 0.049000
Error rate for c= 2820: 490/10000 = 0.049000
Error rate for c= 2830: 490/10000 = 0.049000
Error rate for c= 2840: 490/10000 = 0.049000
Error rate for c= 2850: 488/10000 = 0.048800
Error rate for c= 2860: 487/10000 = 0.048700
Error rate for c= 2870: 488/10000 = 0.048800
Error rate for c= 2880: 487/10000 = 0.048700
Error rate for c= 2890: 487/10000 = 0.048700
Error rate for c= 2900: 488/10000 = 0.048800
Error rate for c= 2910: 488/10000 = 0.048800
Error rate for c= 2920: 488/10000 = 0.048800
Error rate for c= 2930: 488/10000 = 0.048800
Error rate for c= 2940: 489/10000 = 0.048900
Error rate for c= 2950: 489/10000 = 0.048900
Error rate for c= 2960: 488/10000 = 0.048800
Error rate for c= 2970: 490/10000 = 0.049000
Error rate for c= 2980: 490/10000 = 0.049000
Error rate for c= 2990: 490/10000 = 0.049000

In [25]: plt.plot(smoothing_c, error_rates)
plt.xlabel('c (smoothing constant)')
plt.ylabel('error rates')
plt.title('Error rate for varying smoothing constants for 2500 < c < 3400')

Out[25]: <matplotlib.text.Text at 0x1128fc88>

```



2.4.3 Using the validation data set, a couple of candidates for the optimal smoothing_c were found:

Error rate for c= 2860: $487/10000 = 0.048700$

Error rate for c= 2880: $487/10000 = 0.048700$

Error rate for c= 2890: $487/10000 = 0.048700$

2.4.4 Smoothing_c = 2860 was chosen, which yielded an error rate of 4.87% using the validation data set.

2.5 Part (e)

Turn in an iPython that includes:

All your code

Error rate on the MNIST test set

Out of the misclassified test digits, pick 5 at random and display them. For each i

```
In [26]: sc = 2860
         posteriors = []
         for c in classes:
             grouping = class_grouping(c)
             mean = np.array(grouping.mean(0))[0]
             cov = np.cov(grouping, rowvar=0)
```

```

cov_smoothed = cov + (sc * np.eye(mean.shape[0]))
Px = multivariate_normal(mean, cov_smoothed)
posteriors.append(Px)

Y = []
for x in X_test:
    bayes_prob = []
    for c in classes:
        prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
        bayes_prob.append(prob)
    prediction = max(bayes_prob, key= lambda a: a[1])
    Y.append(prediction[0])

errors = (y_test != Y).sum()
total = X_test.shape[0]
error_rate = errors/float(total)
print("Error rate for c= %s: %d/%d = %f" % (sc, errors, total, error_rate))

```

Error rate for c= 2860: 436/10000 = 0.043600

2.6 The error rate on the MNIST test set is: 4.36%

```


In [27]: print 'The model (no validation set used) has an accuracy of', metrics.accuracy_score(y_test, Y)
print metrics.classification_report(y_test, Y)

cm = pd.DataFrame(metrics.confusion_matrix(y_test, Y))
plt.figure(figsize=(15, 15))
sns.heatmap(cm, annot=True, fmt=".3f", linewidth=.5, square = True, cmap = 'magma')
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix', size = 15);

```

The model (no validation set used) has an accuracy of 0.9564

	precision	recall	f1-score	support
0	0.97	0.99	0.98	980
1	0.96	0.99	0.97	1135
2	0.97	0.94	0.95	1032
3	0.95	0.94	0.95	1010
4	0.98	0.96	0.97	982
5	0.97	0.94	0.95	892
6	0.97	0.97	0.97	958
7	0.97	0.93	0.95	1028
8	0.91	0.94	0.92	974
9	0.92	0.96	0.94	1009
avg / total	0.96	0.96	0.96	10000



2.6.1 The model was able to correctly predict the majority of the diagonals, with an accuracy of 95.64%. Digits 8 & 9 had the lowest precision (but still > 0.90).

```
In [28]: indices = np.array(np.where((y_test != Y)==True))[0]

        wrong = random.sample(indices, 5)
        actuals = y_test[wrong]
        misclassified_predictions = []
        for x in X_test[wrong]:
```

```

bayes_prob = []
for c in classes:
    prob = [c, np.log(priors[c]) + posteriors[c].logpdf(x)]
    bayes_prob.append(prob)
prediction = max(bayes_prob, key= lambda a: a[1])
misclassified_predictions.append(prediction[0])

for i in range(len(wrong)):
    bayes_prob = []
    for c in classes:
        prob = [c, np.log(priors[c]) + posteriors[c].logpdf(X_test[wrong[i]])]
        bayes_prob.append(prob)

    print 'The Bayes probability found is \n', bayes_prob
    print 'For this random example the model predicted a \n', misclassified_predictions[i]
    print ("Let's see how it compares to the actual image: {}".format(actual_image[wrong[i]]))
    plt.figure(1, figsize=(3, 3))
    plt.imshow(X_test[wrong[i]].reshape(28, 28), cmap=plt.cm.gray_r, interpolation='nearest')
    plt.show()

```


The Bayes probability found is

```
[[0, -4086.2044056835412], [1, -4035.5495631430299], [2, -4071.9587276439979], [3, -4035.5495631430299]]
```

For this random example the model predicted a

8

Let's see how it compares to the actual image: 5




The Bayes probability found is

```
[[0, -4102.4500641906534], [1, -4151.553061514458], [2, -4033.8034182979213], [3, -4033.8034182979213]]
```

For this random example the model predicted a

8

Let's see how it compares to the actual image: 2




The Bayes probability found is

```
[[0, -4174.8470824229125], [1, -4118.5923467660514], [2, -4117.8467729428721], [3,
```

For this random example the model predicted a

9


Let's see how it compares to the actual image: 4



The Bayes probability found is

```
[[0, -4132.2417508921853], [1, -4173.2826738035337], [2, -4049.3830047133774], [3,  
For this random example the model predicted a  
7
```

Let's see how it compares to the actual image: 2



The Bayes probability found is

```
[[0, -4064.0727425332047], [1, -4355.1813880862801], [2, -4068.8853861709608], [3,  
For this random example the model predicted a  
6
```

Let's see how it compares to the actual image: 4

